

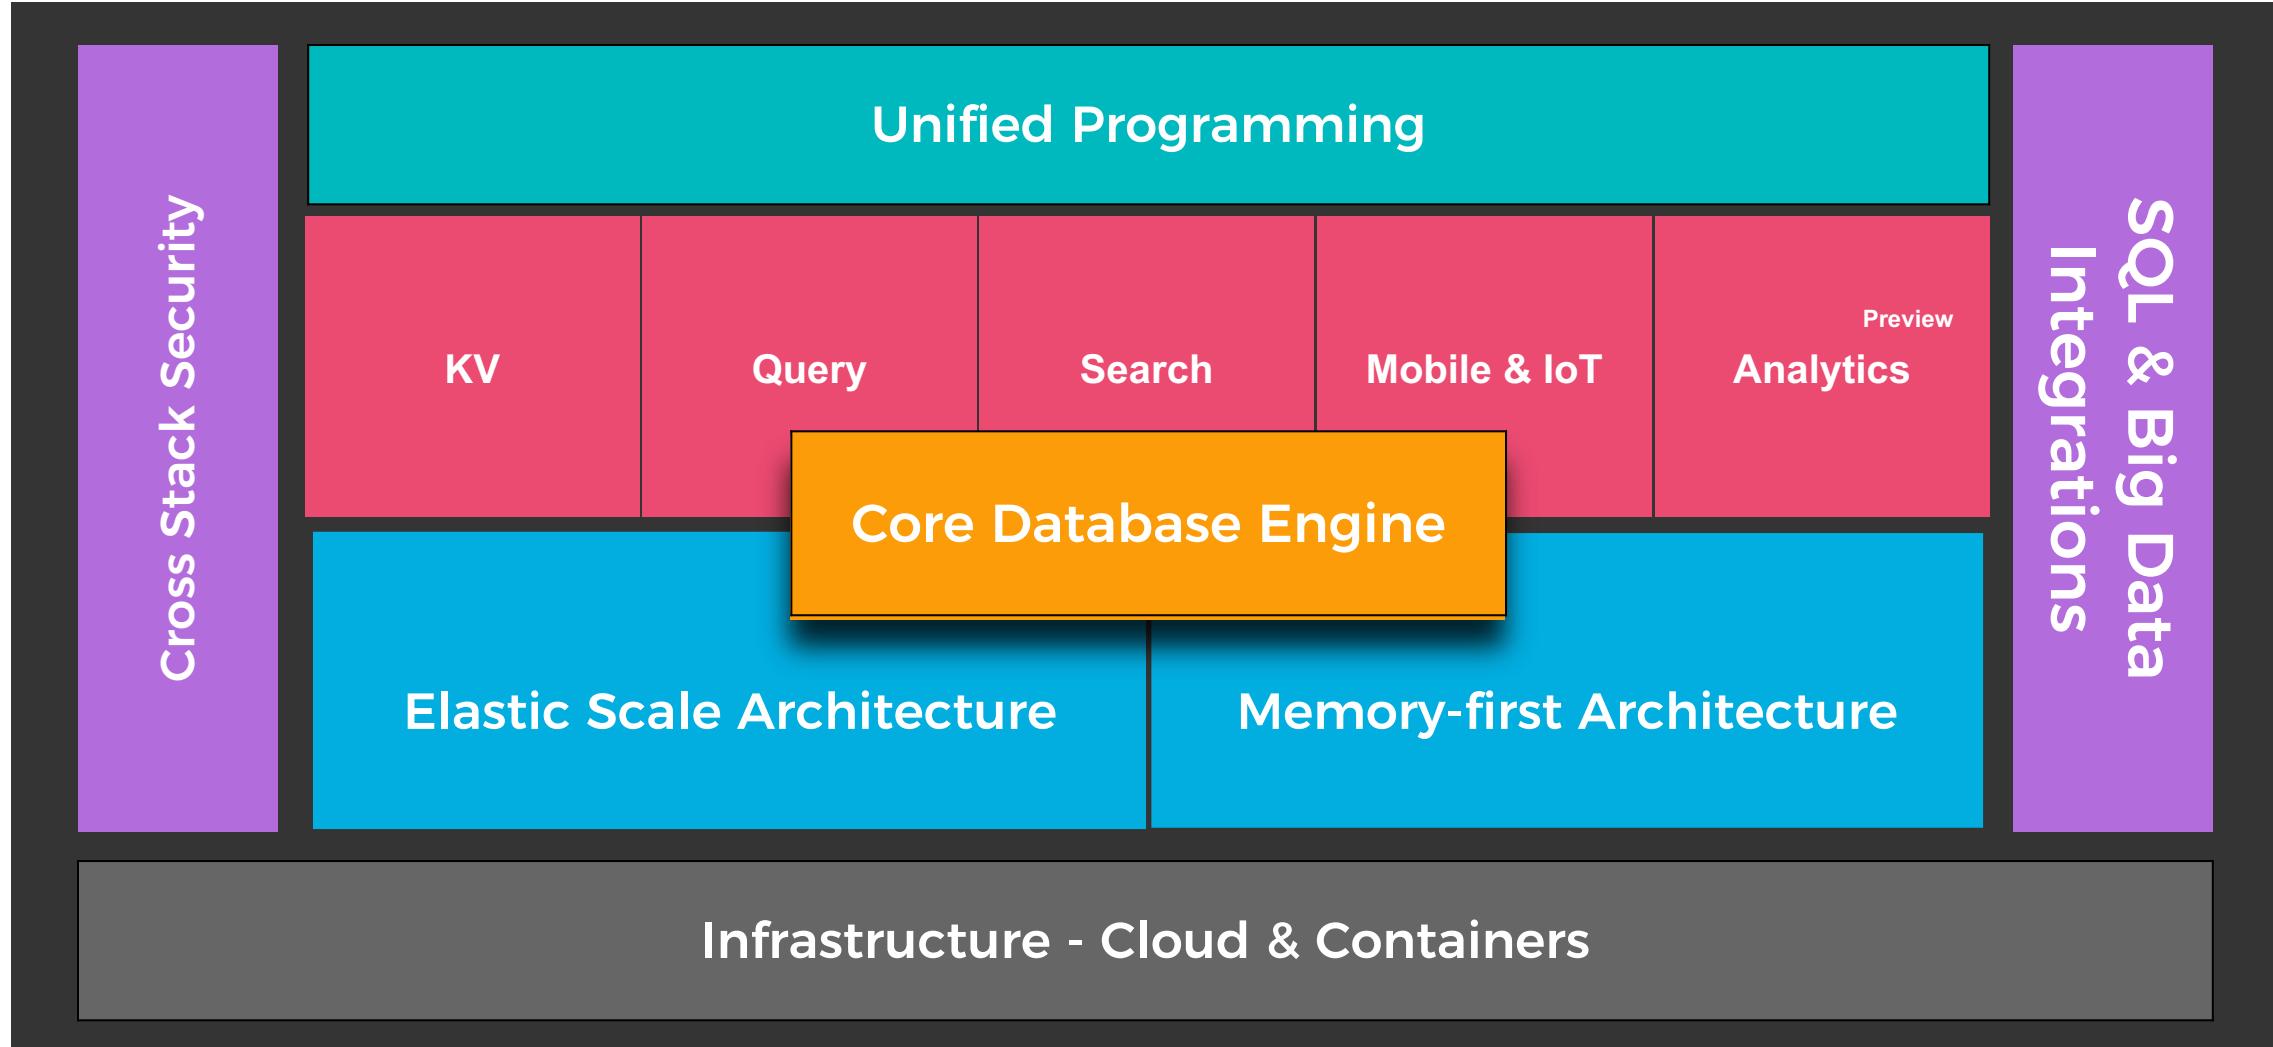


Couchbase

Architecture and Administration Basics

Workshop Day 1 - Architecture

Couchbase Data Platform



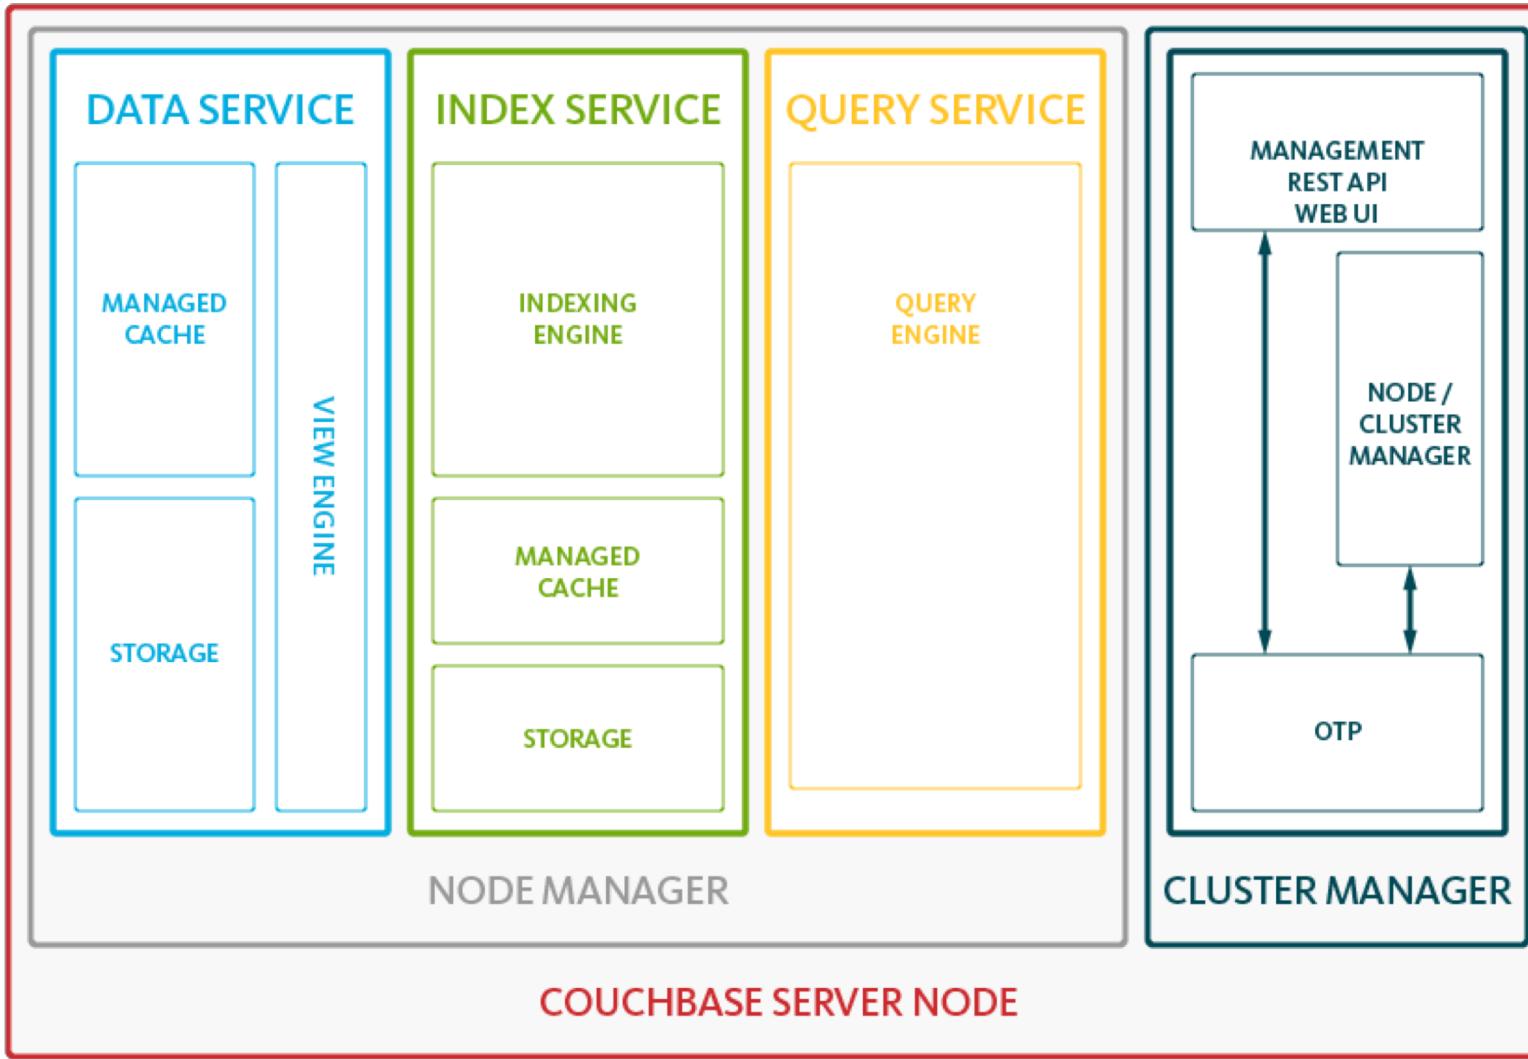


1

Architecture



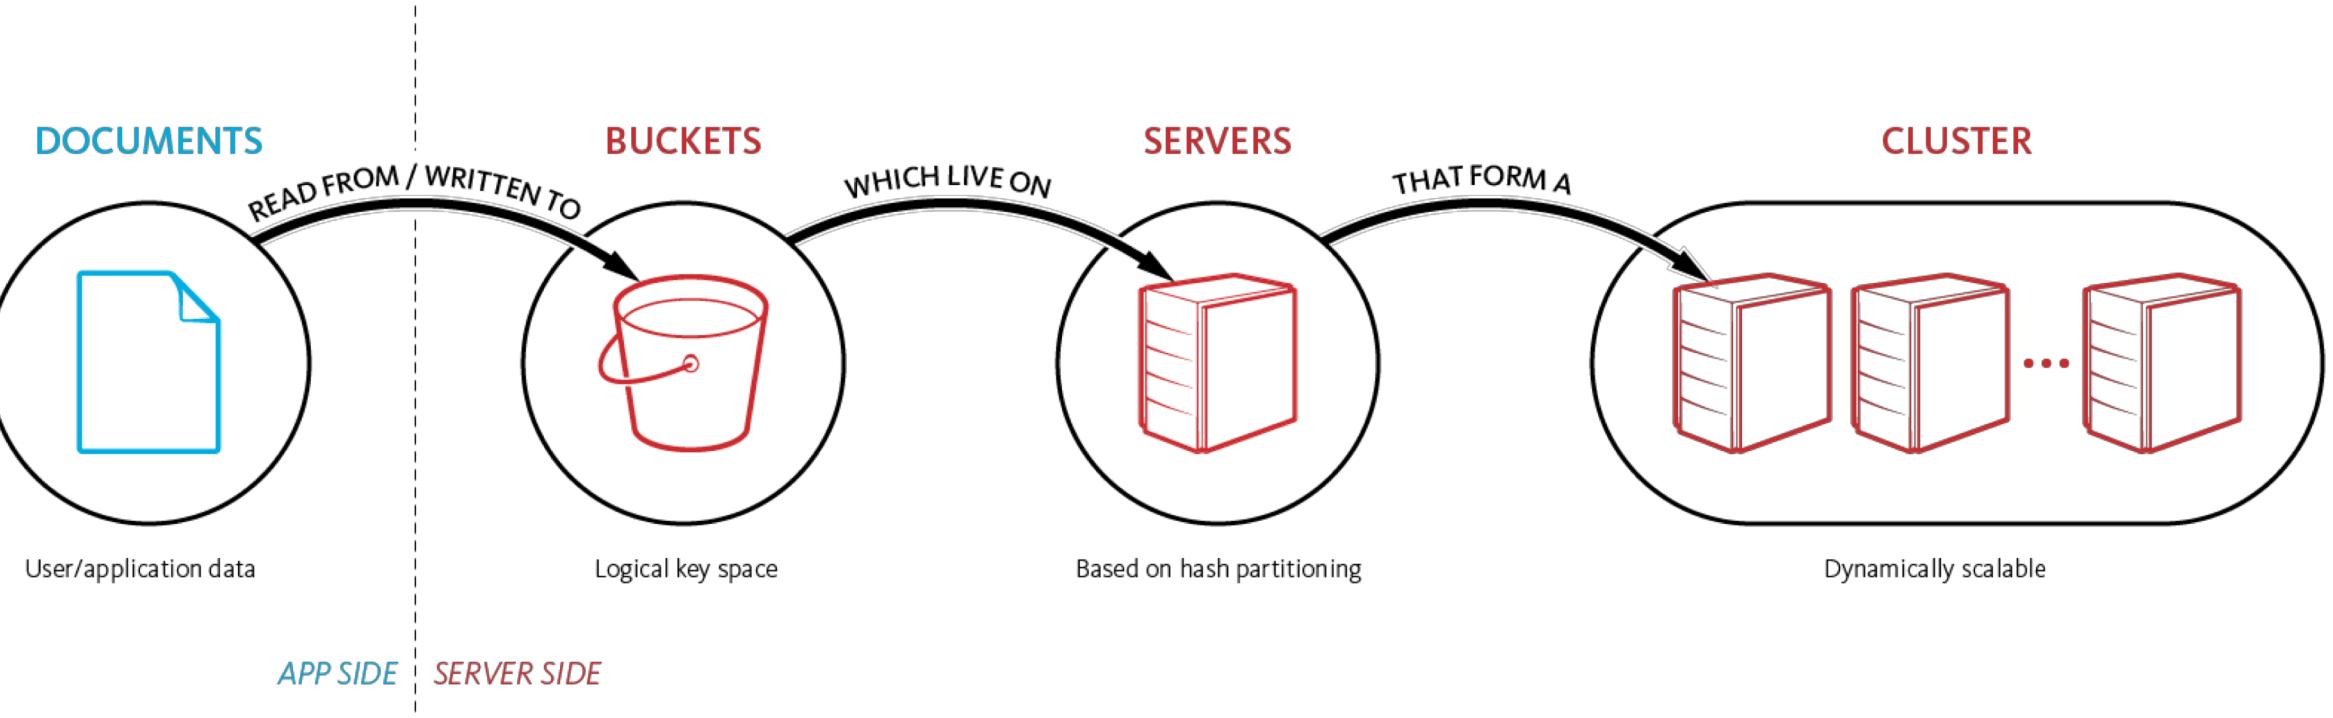
Couchbase Architecture



- **Data Service** - Key Value Store and builds and maintains Distributed secondary indexes (MapReduce Views)
- **Indexing Engine** - builds and maintains Global Secondary Indexes
- **Query Engine** - plans, coordinates, and executes queries against either Global or Distributed indexes
- **Cluster Manager** - configuration, heartbeat, statistics, RESTful Management interface



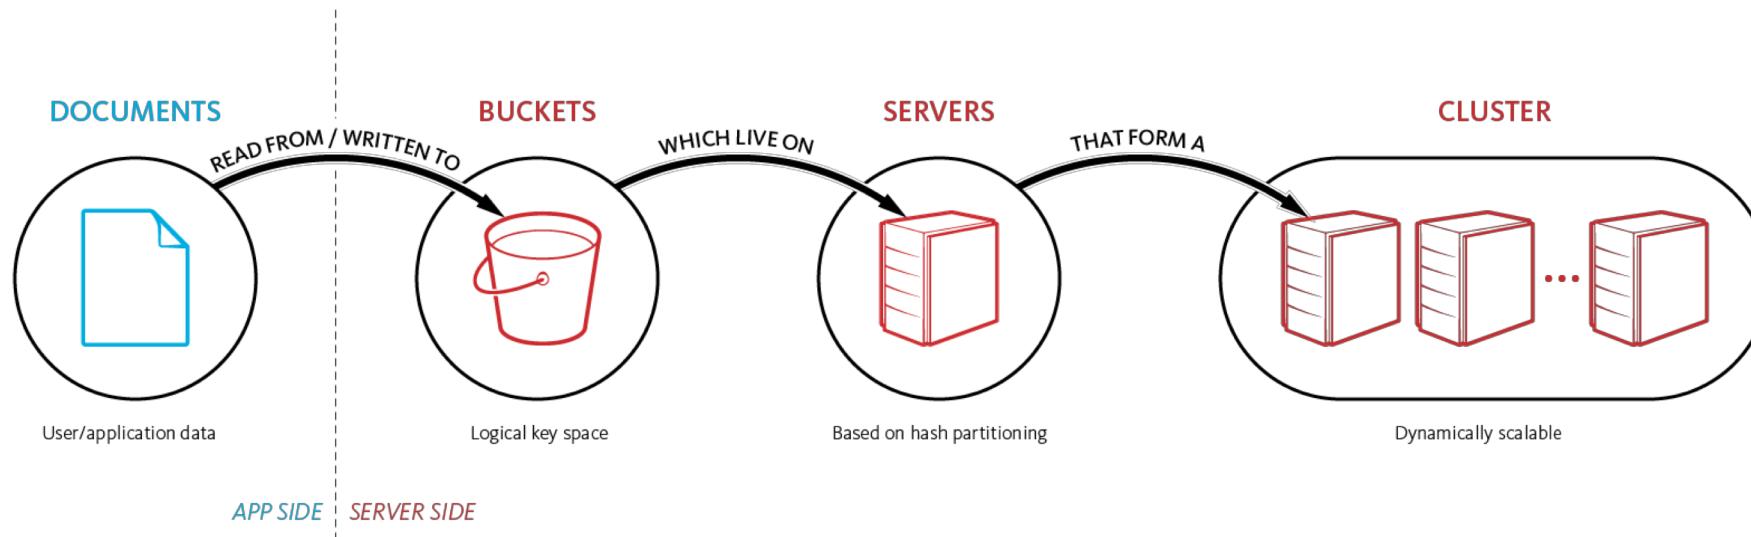
Storing And Retrieving Documents





Buckets - When to use more than one ?

- When you need to treat or access the data differently
 - Different High Availability requirements (1,2 or 3 replicas)
 - Different performance / residency needs (how much data to cache)
 - Security / Multi-tenancy
 - Segregating Binary and JSON data - especially with view usage





Auto sharding - Bucket and vBuckets



Virtual buckets

- Bucket
 - A **bucket** is a logical, unique key space
 - Multiple buckets can exist within a single cluster of nodes
- vBuckets
 - Each bucket has active and replica data sets (1, 2 or 3 extra copies)
 - Each data set has **1024 Virtual Buckets** (vBuckets)
 - Each vBucket contains 1/1024th portion of the data set
 - vBuckets do not have a fixed physical server location
 - Mapping between the vBuckets and physical servers is called the **cluster map**
 - Document IDs (keys) always get hashed to the same vBucket (consistent hashing)
 - Couchbase SDK's lookup the vBucket -> server mapping



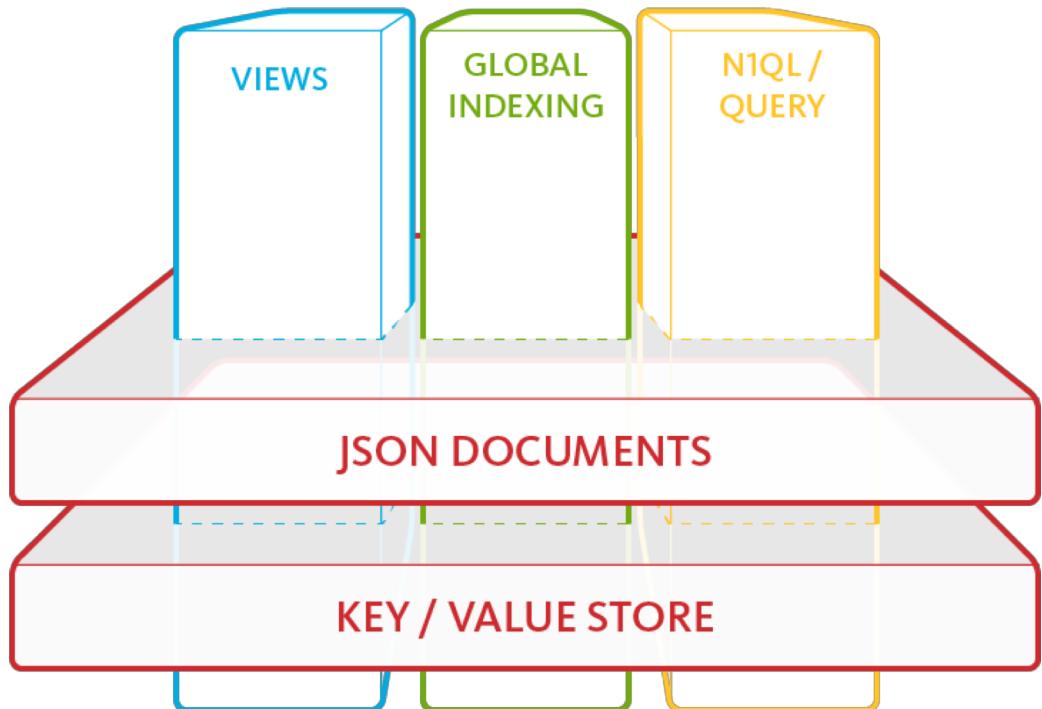
Bucket Comparison

	Memcached	Couchbase	Ephemeral <small>New in 5.0</small>
Persistence	X		X
Replication	X		
Rebalance	X		
XDCR	X		
N1QL / Indexing	X		*
Analytics/ Eventing	X		
Max Object Size	1MB	20MB	20MB

* MOI, FTS only



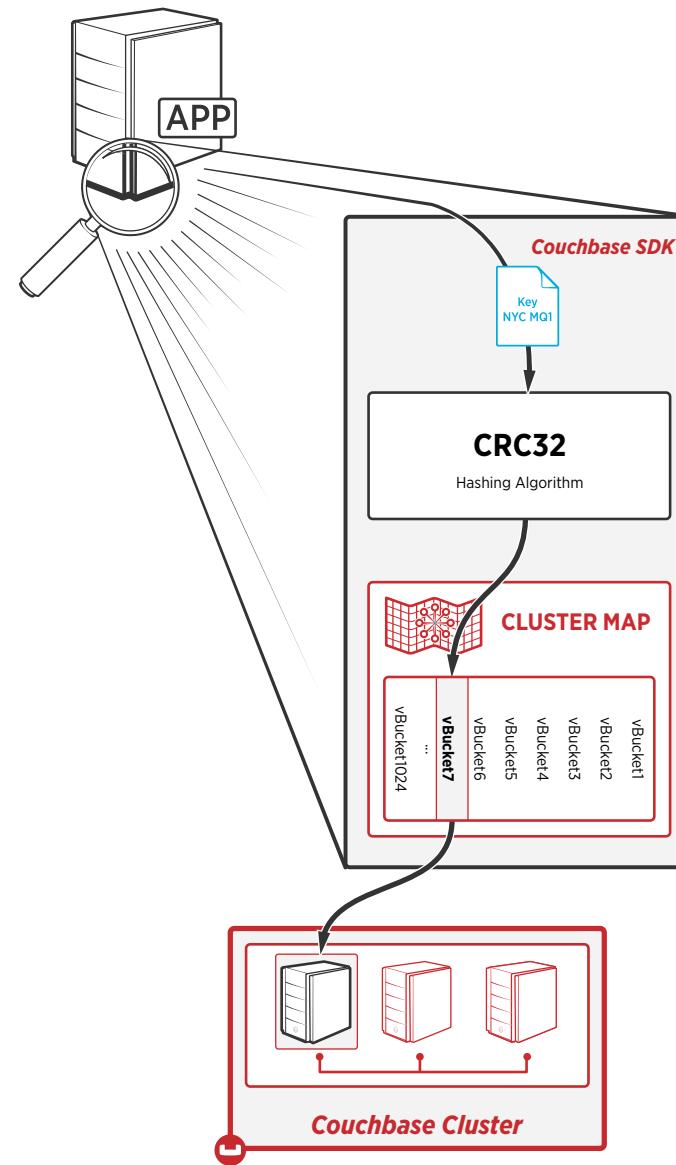
Couchbase Data Access



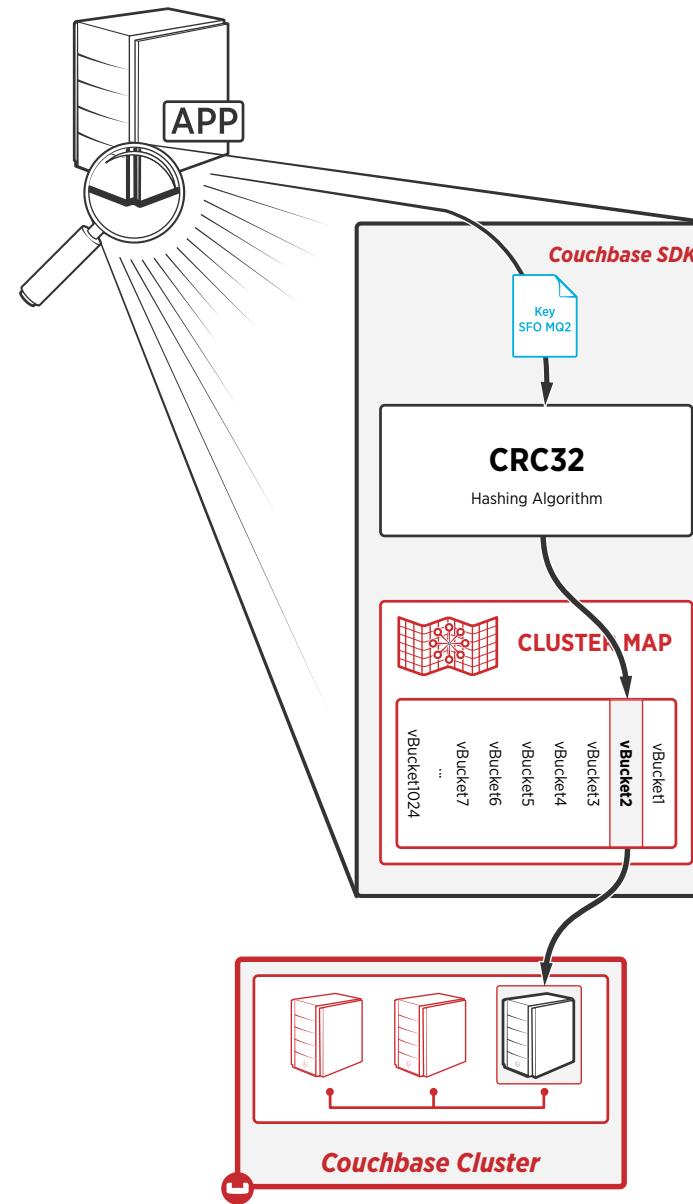
- Everything is layered on top of **Key Value**
- A **Document** store is a special case of Key-Value
- **Views** provide aggregation and real-time analytics through incremental map-reduce
- **Global Secondary Indexes** provide low latency/high throughput indexes
- **N1QL** is a language that provides a powerful and expressive way of accessing documents



Cluster Map

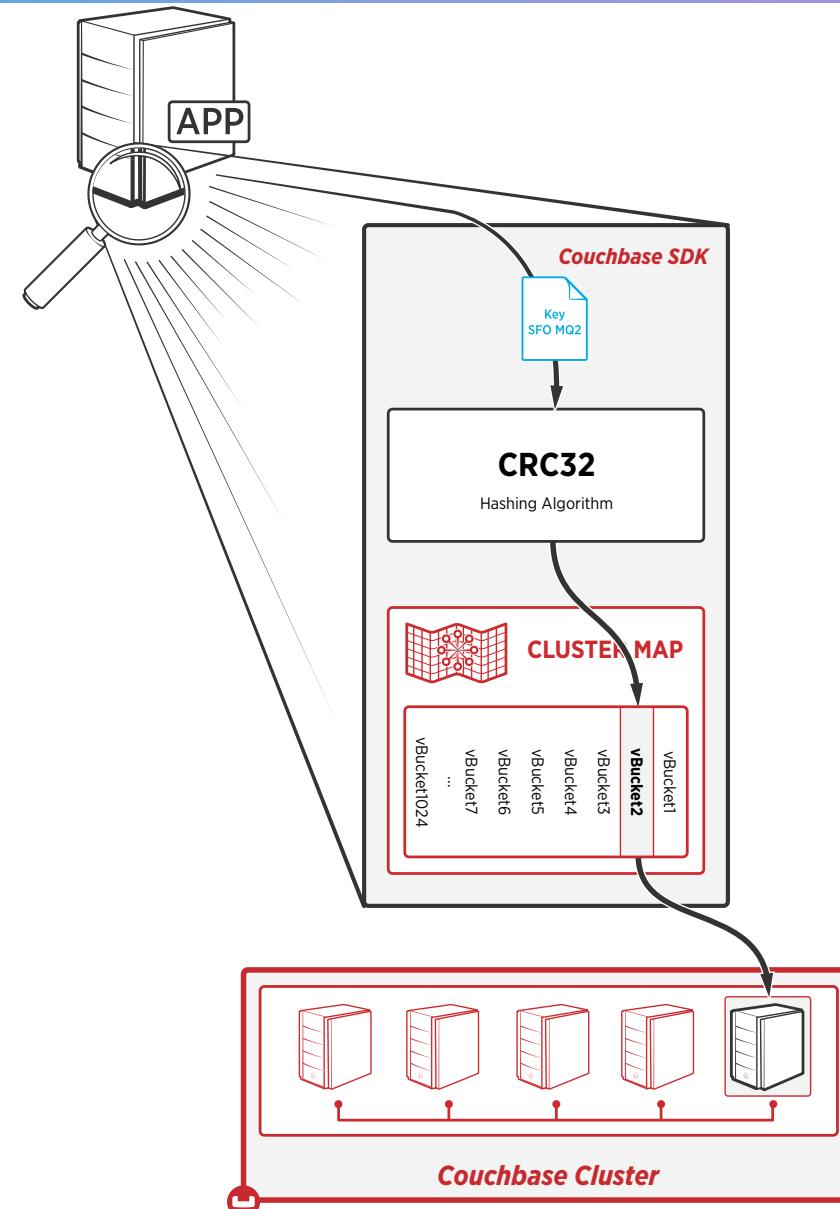


Cluster Map



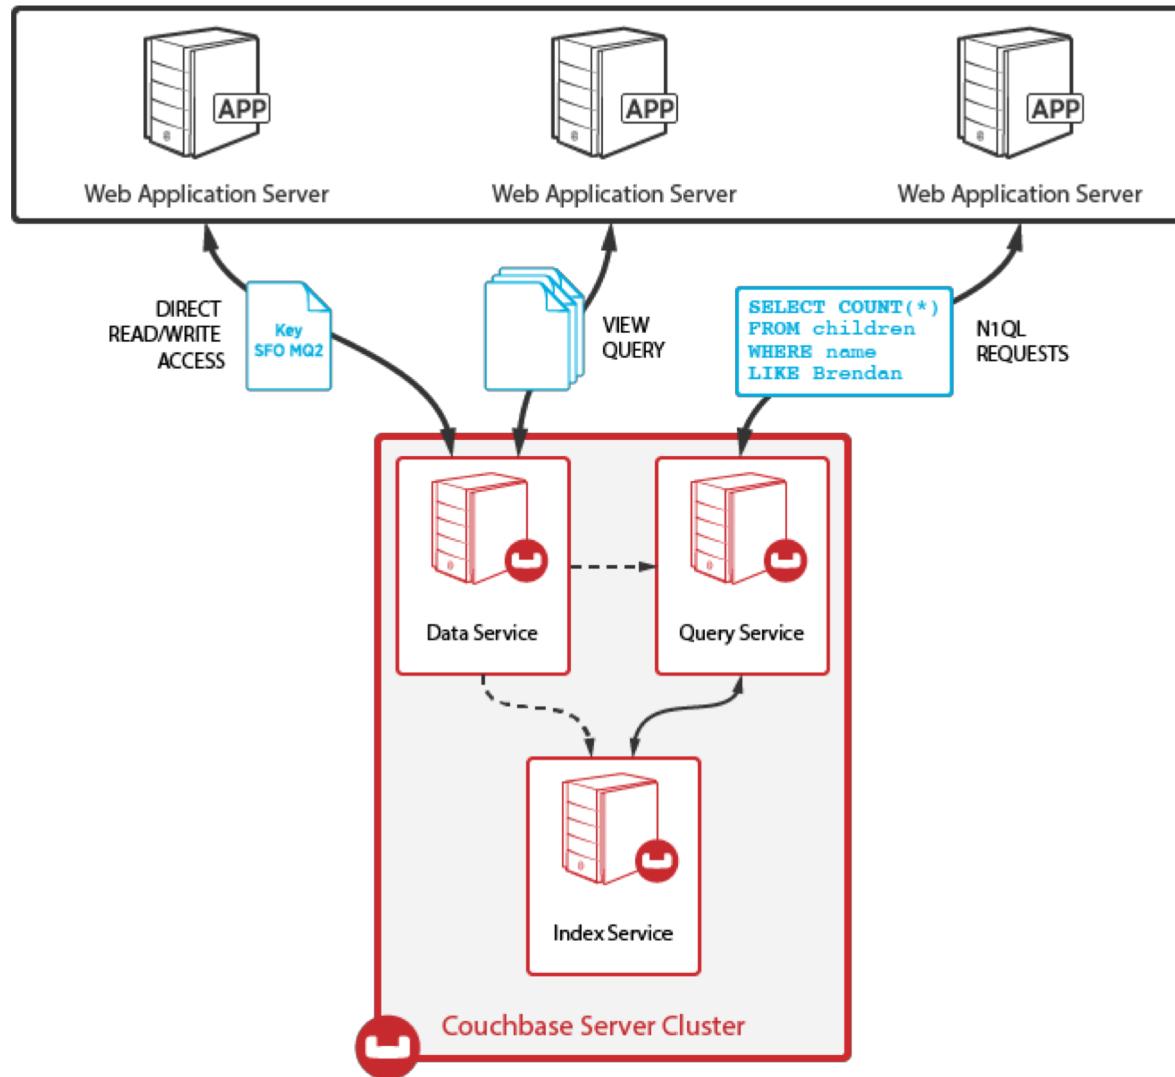


Cluster Map - Addition of 2 Nodes





Application to Database Interaction



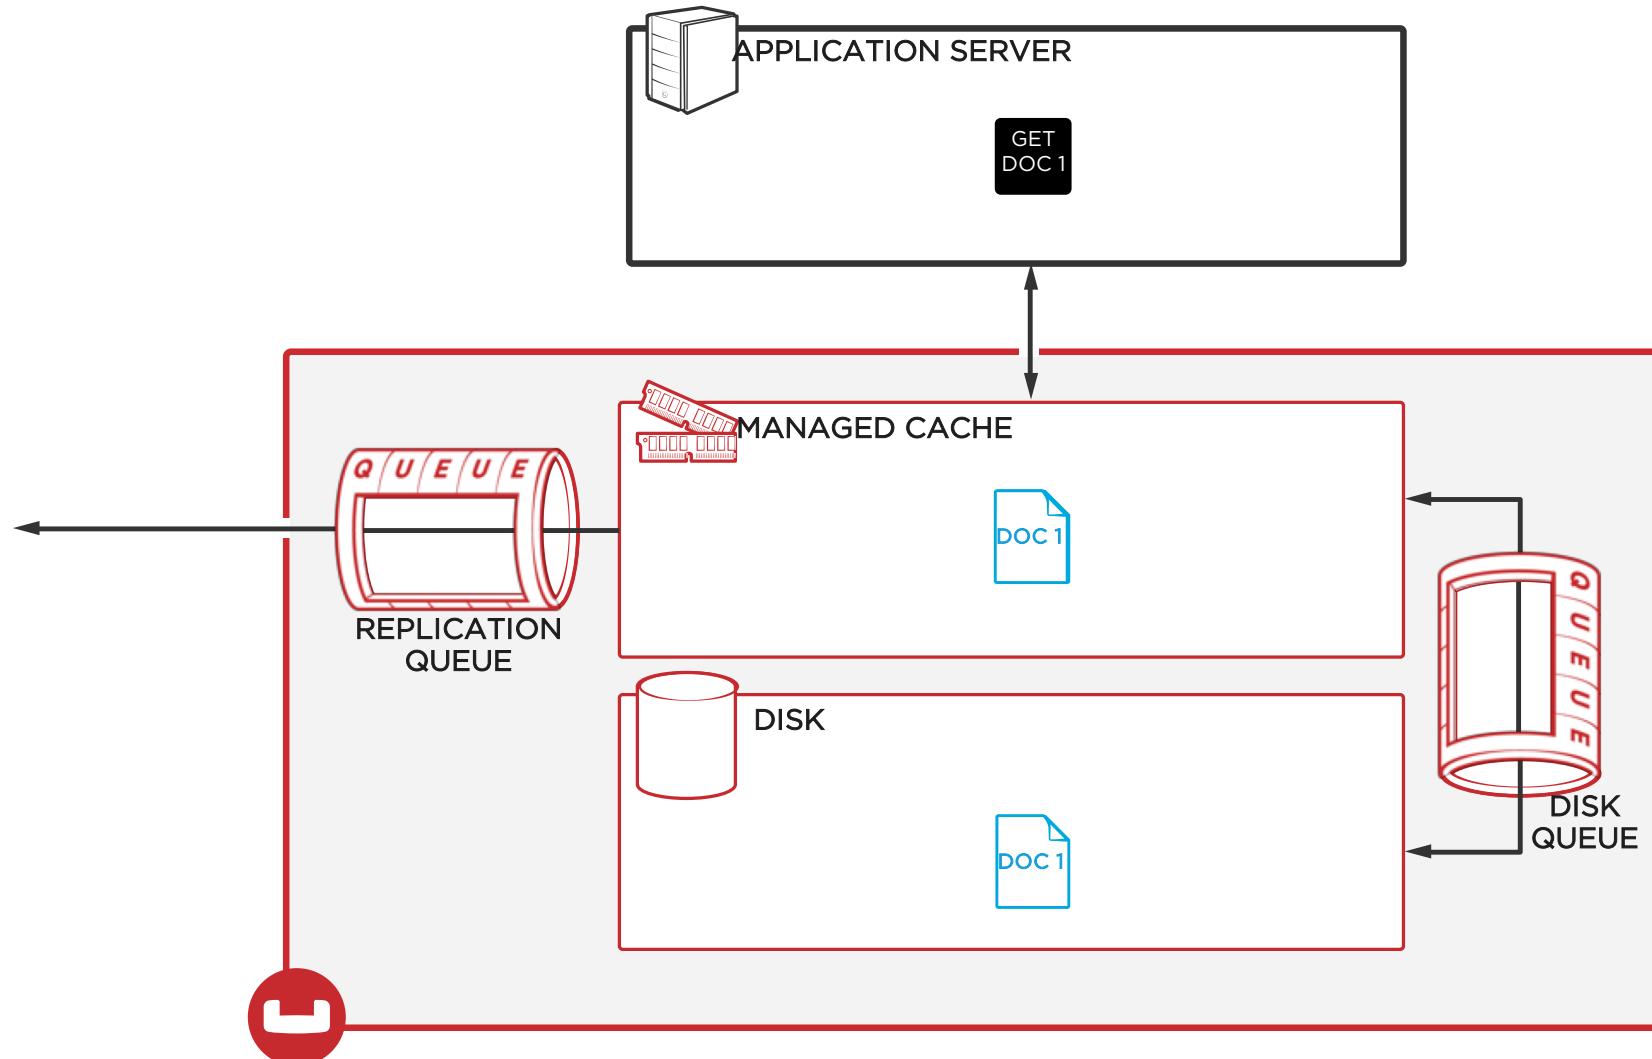


2

SDK Operations

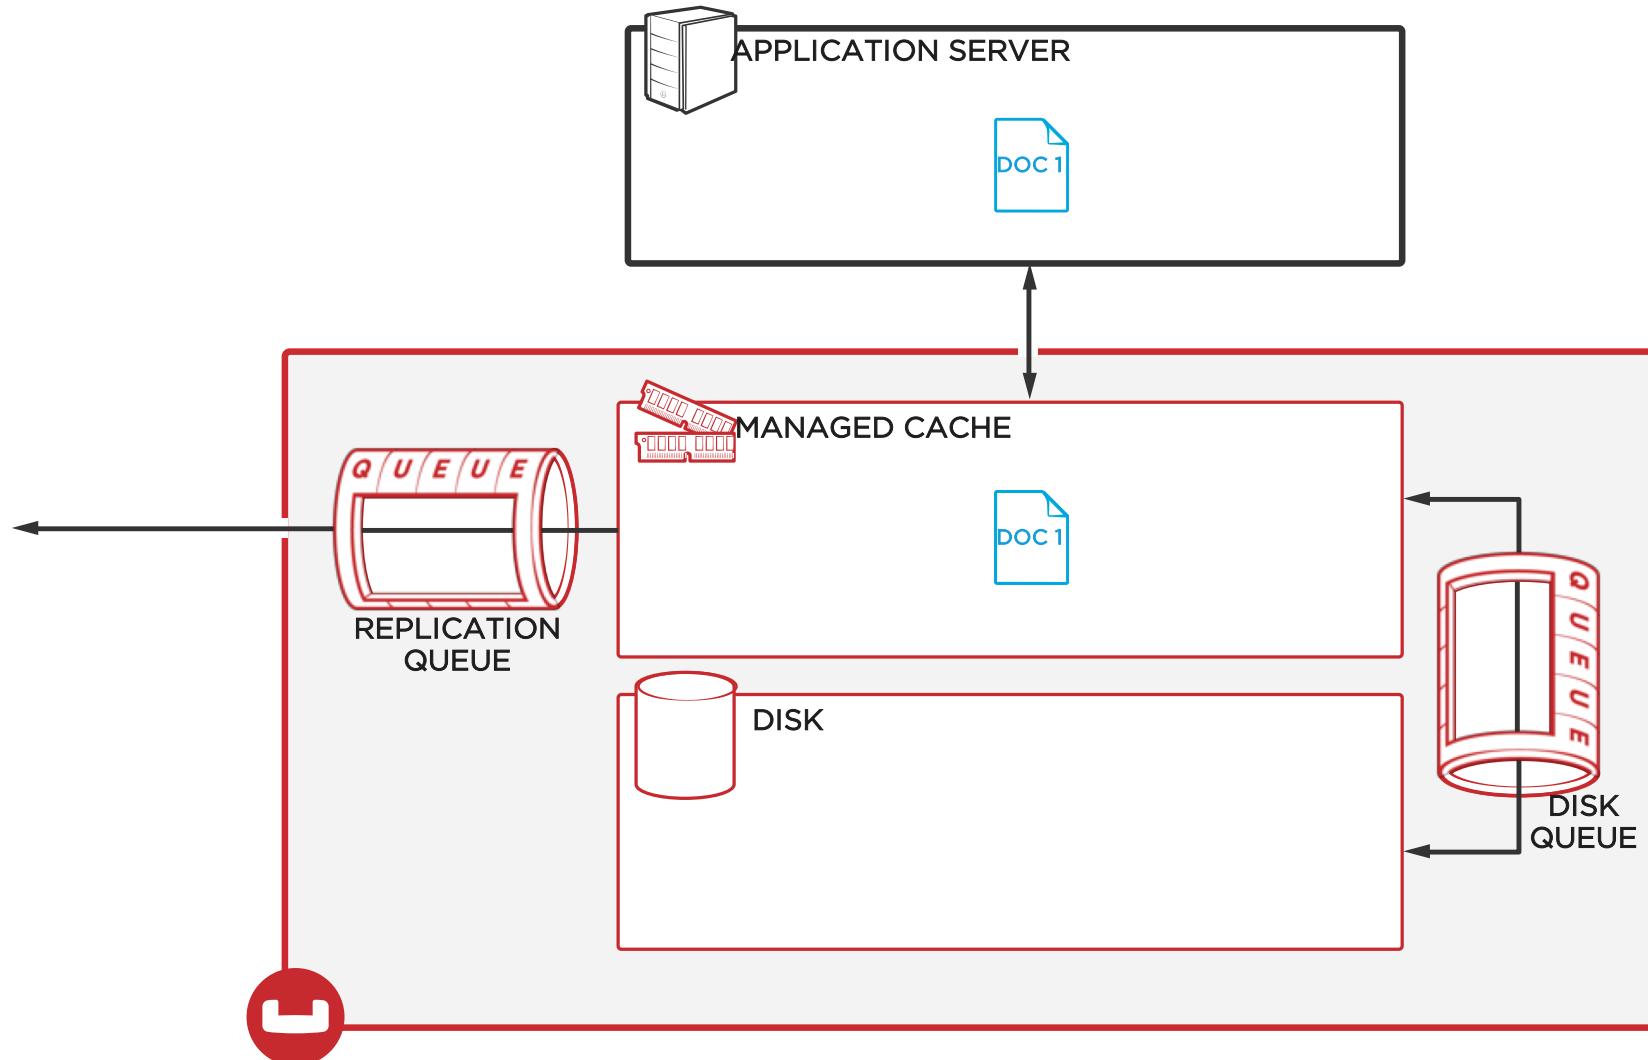


Couchbase Read Operation



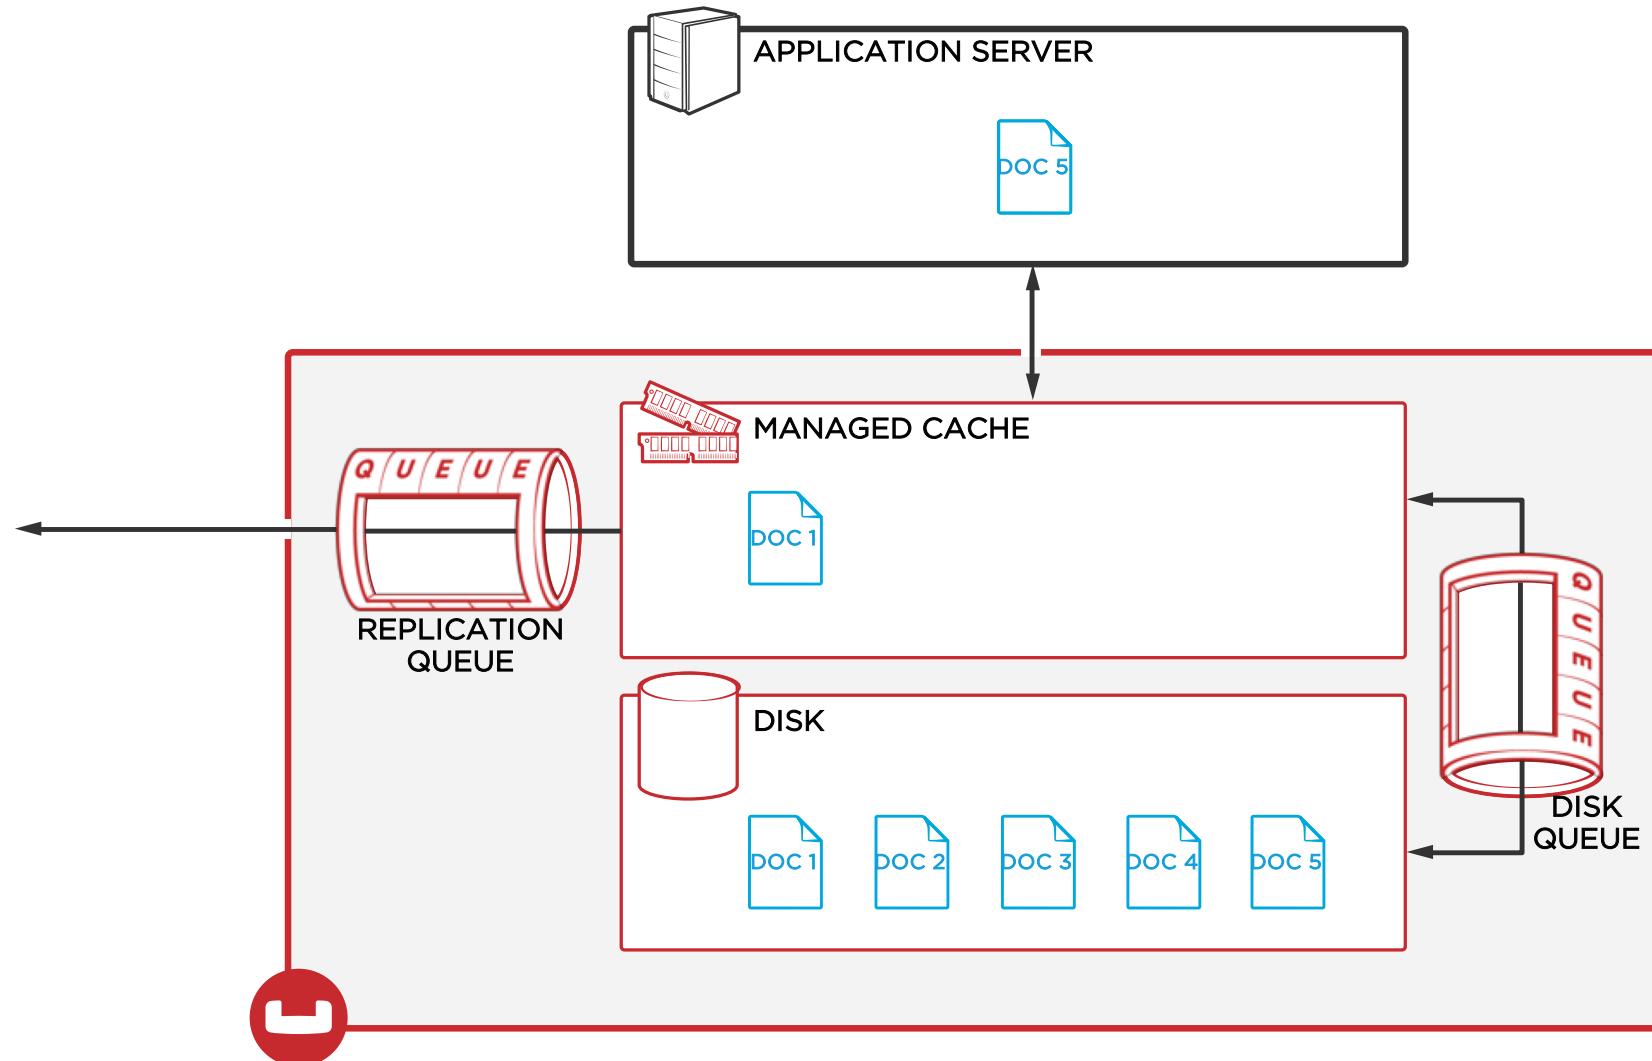


Write Operation



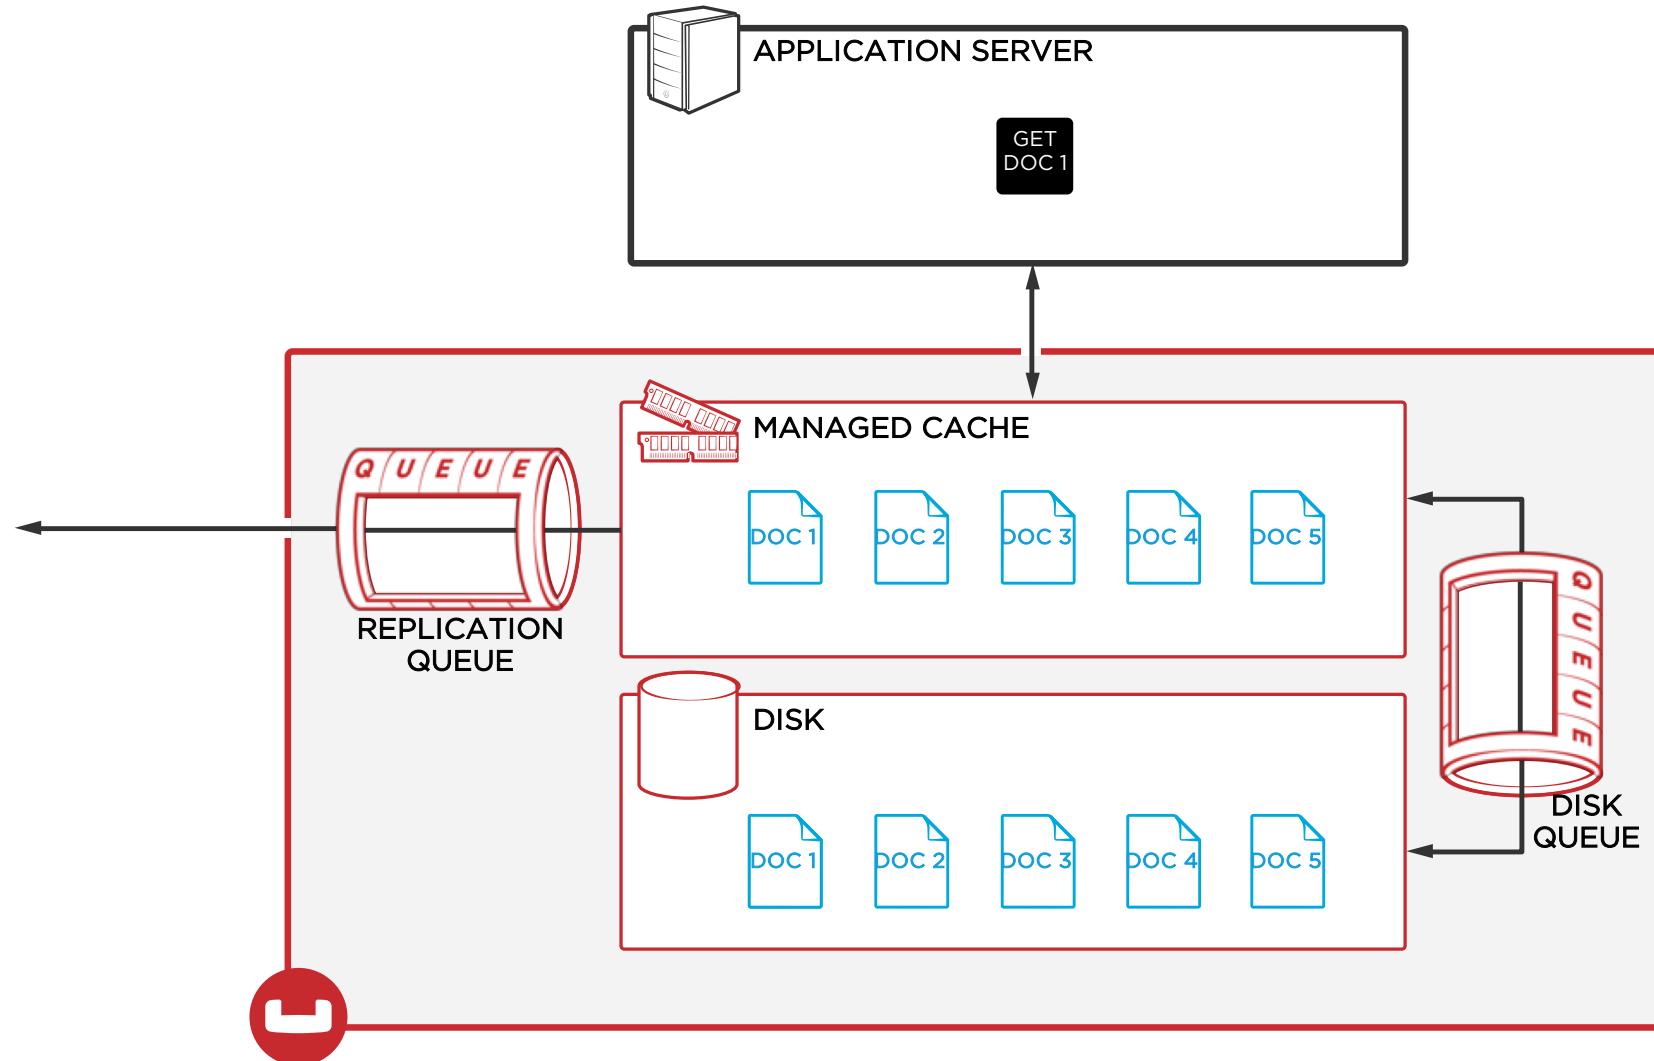


Cache Ejection





Cache Miss



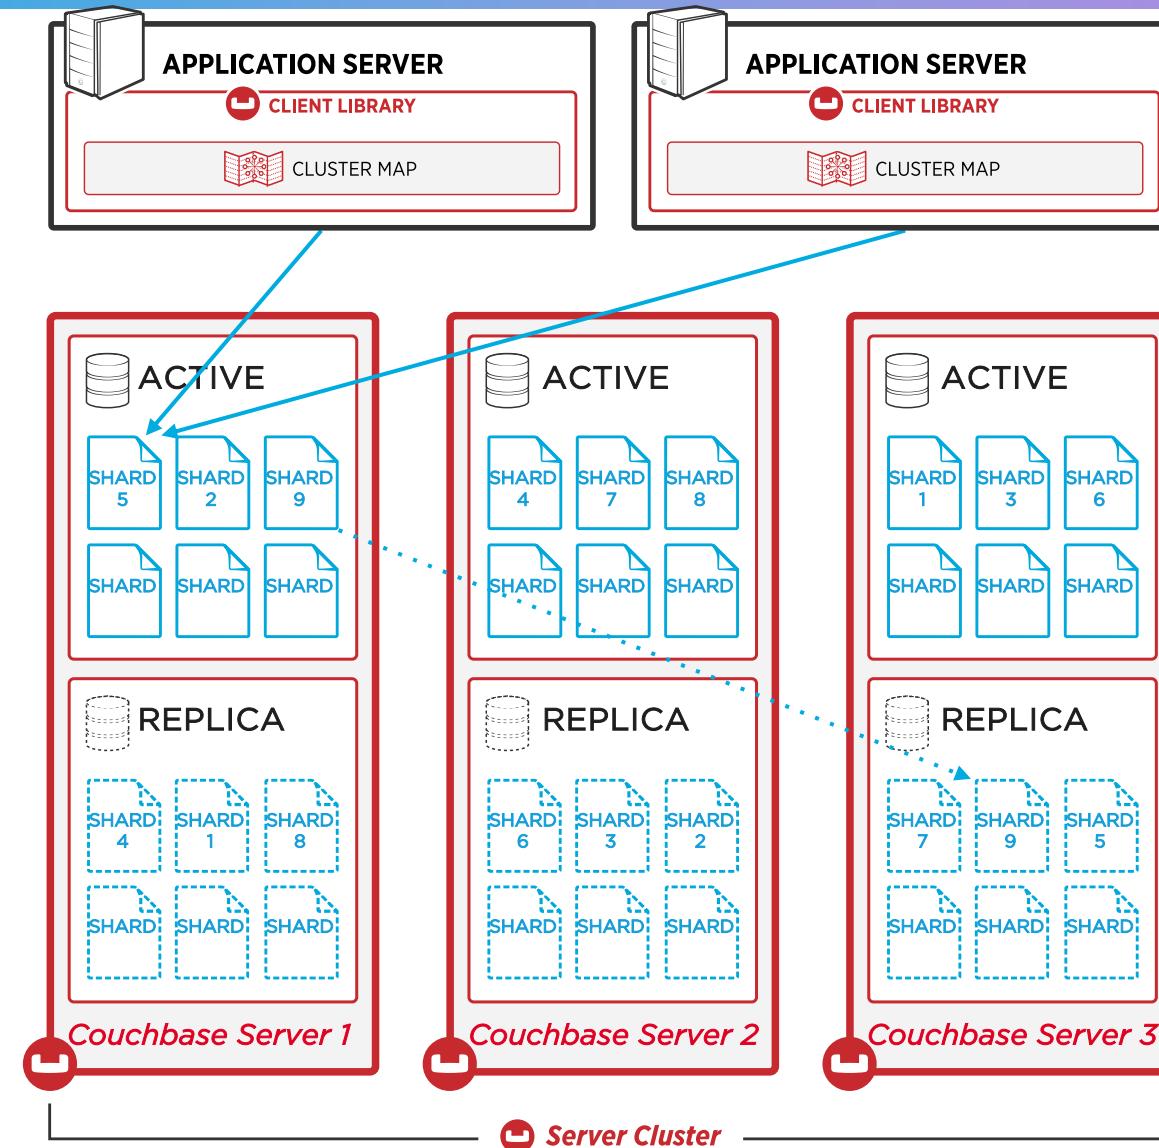


3

Cluster Operations

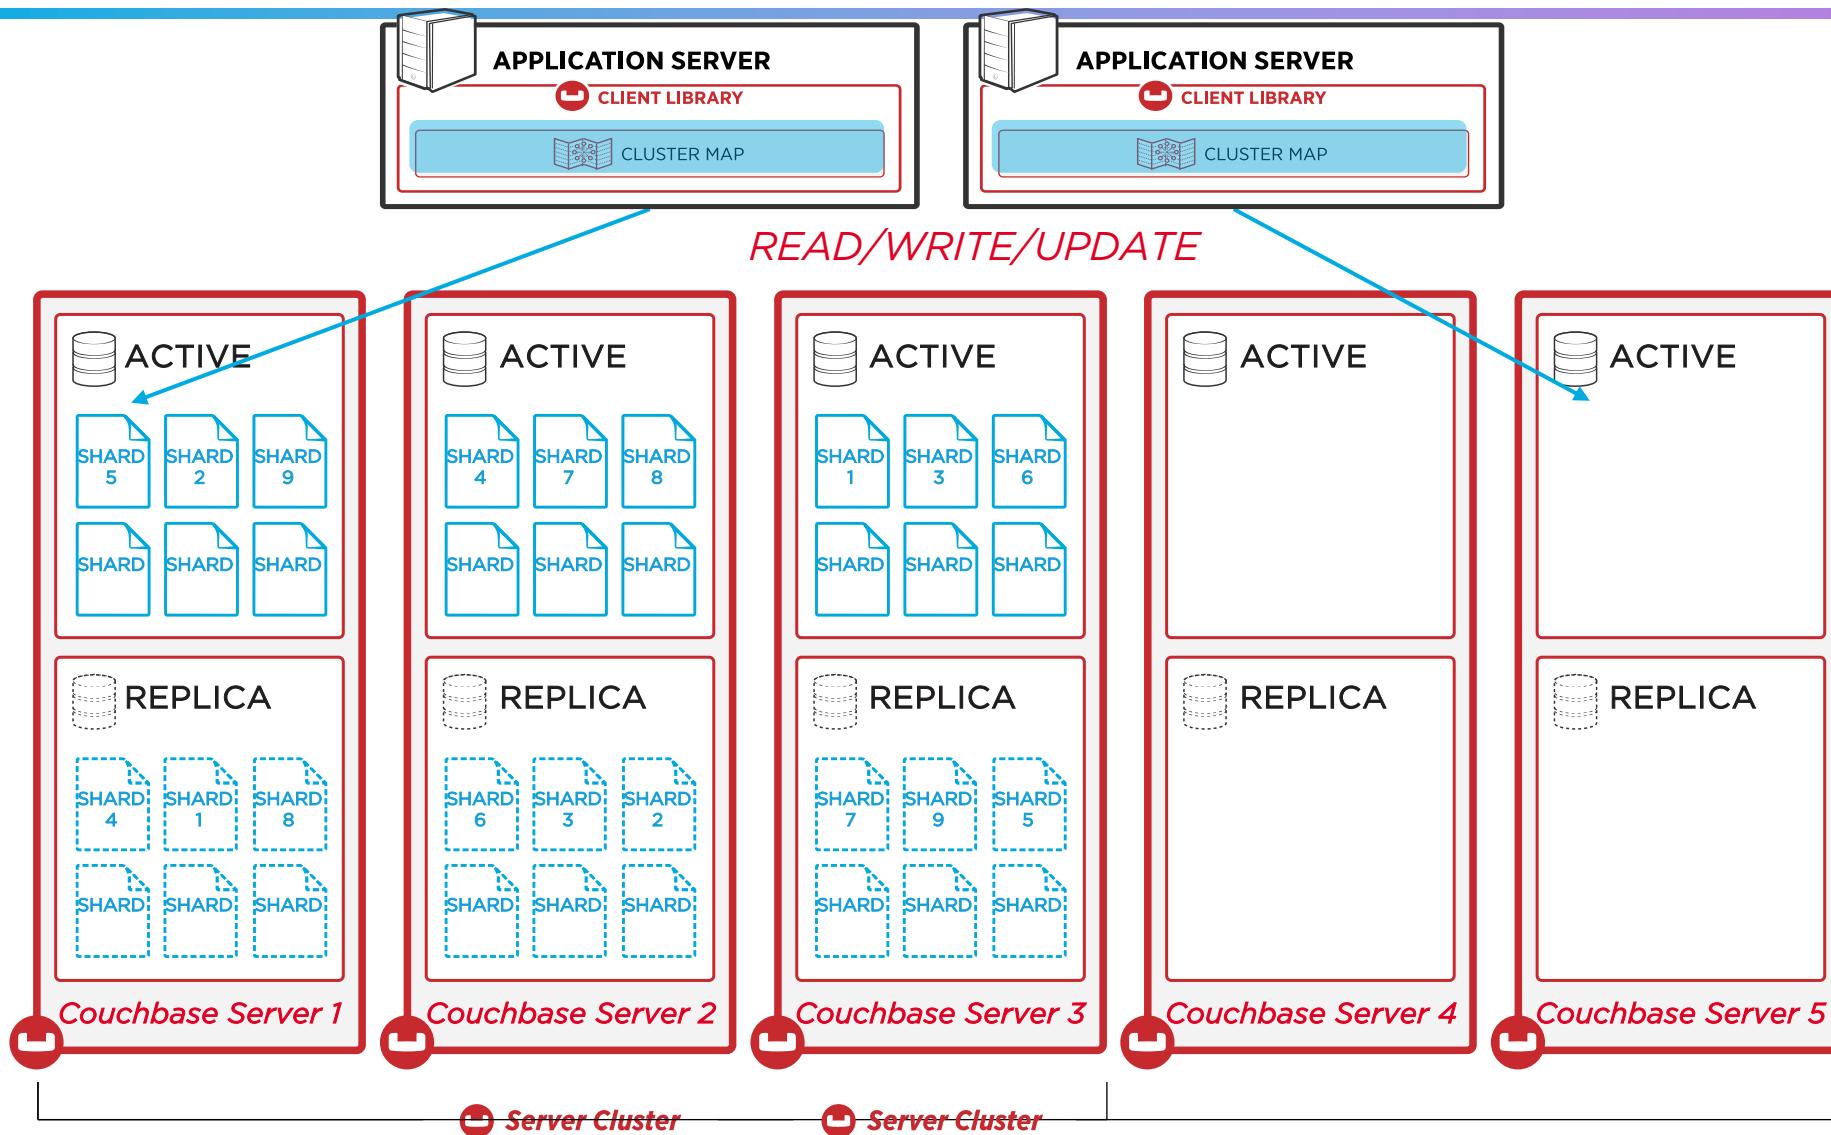


Basic Operation



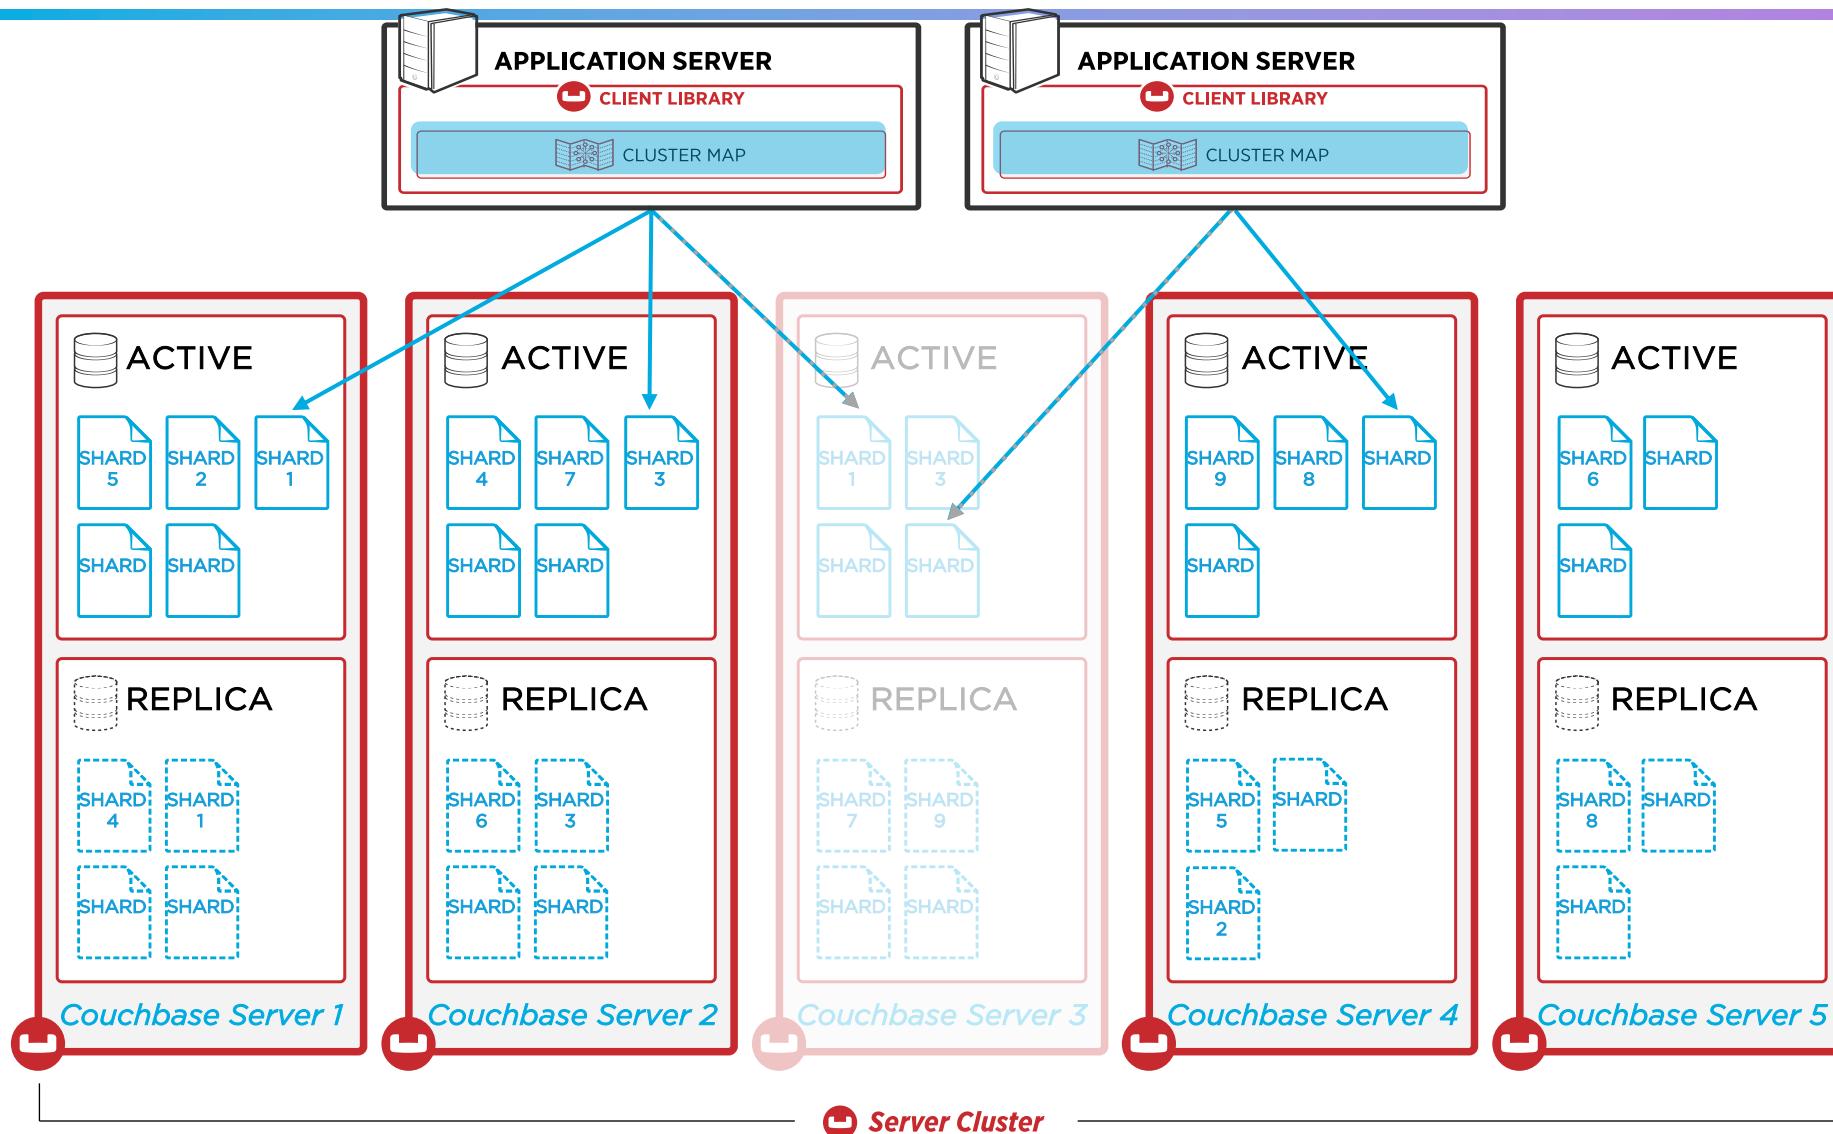


Add Nodes to Cluster





Fail Over Node



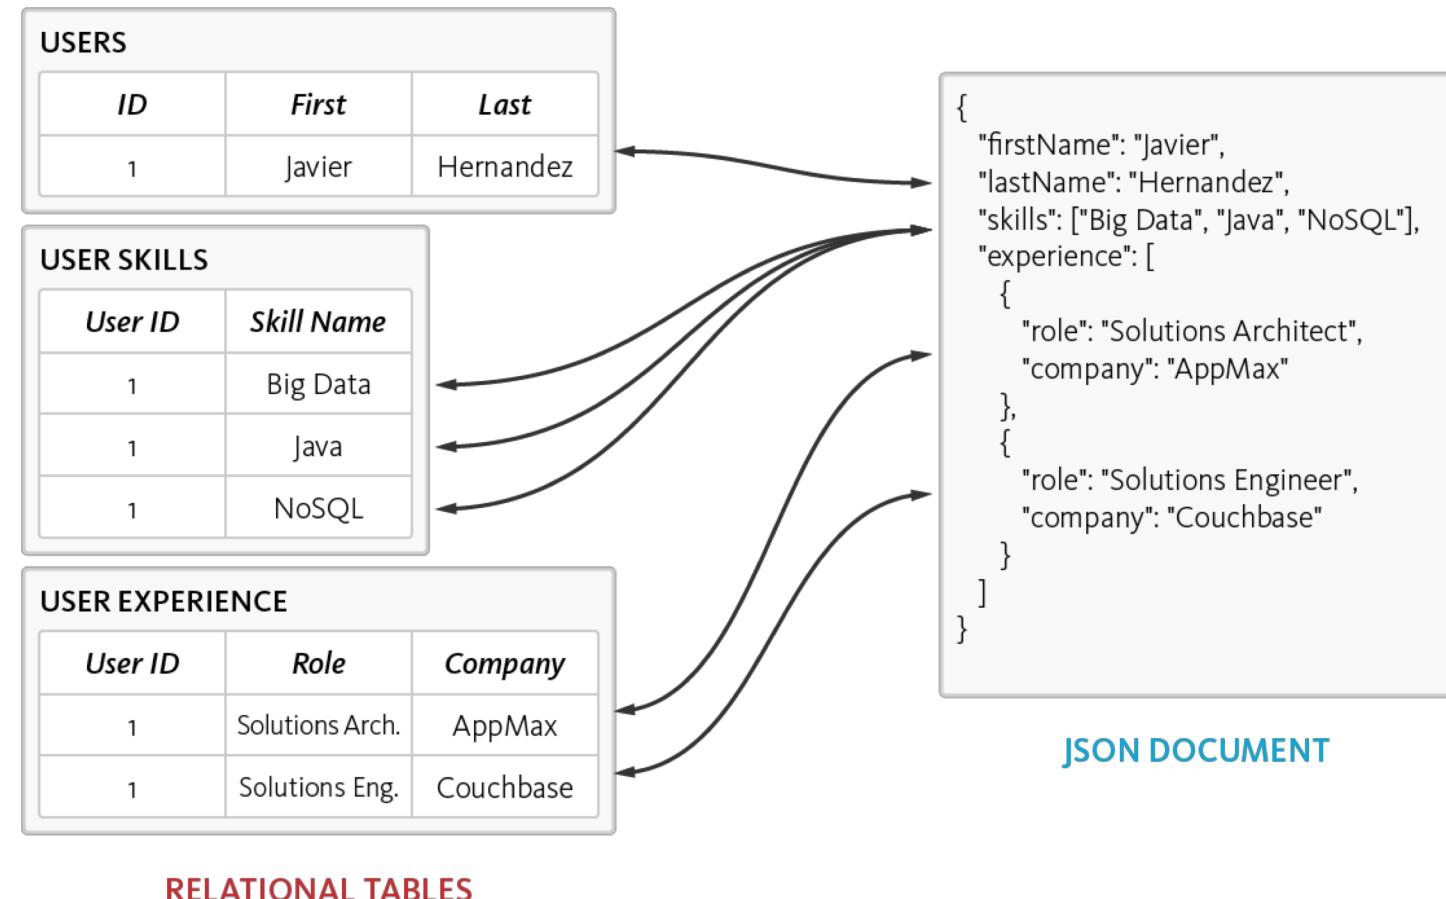
JSON Support

Flexibility



The Power of the Flexible JSON Schema

- Ability to store data in multiple ways
 - Denormalized single document, as opposed to normalizing data across multiple table
 - Dynamic Schema to add new values when needed



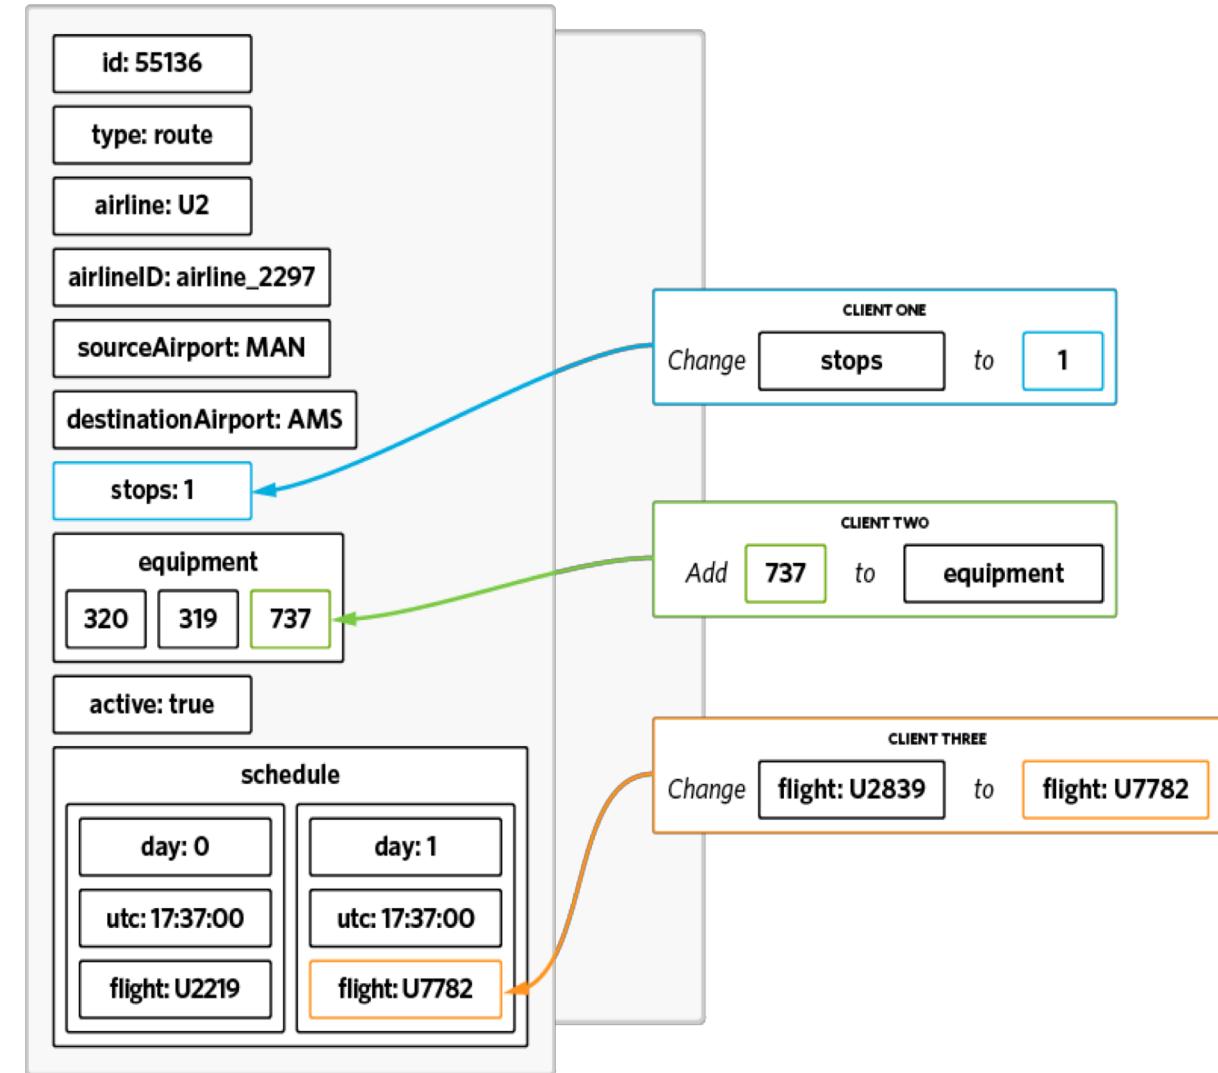


Efficient Sub-Document Operations

- Document Mutations:

- Atomic Operate on individual fields
- Identical syntax behavior to regular bucket methods
(upsert, insert, get, replace)
- Support for JSON fragments.
- Support for Arrays with uniqueness guarantees and ordinal placement (front/back)

Document Key: "route_55136"

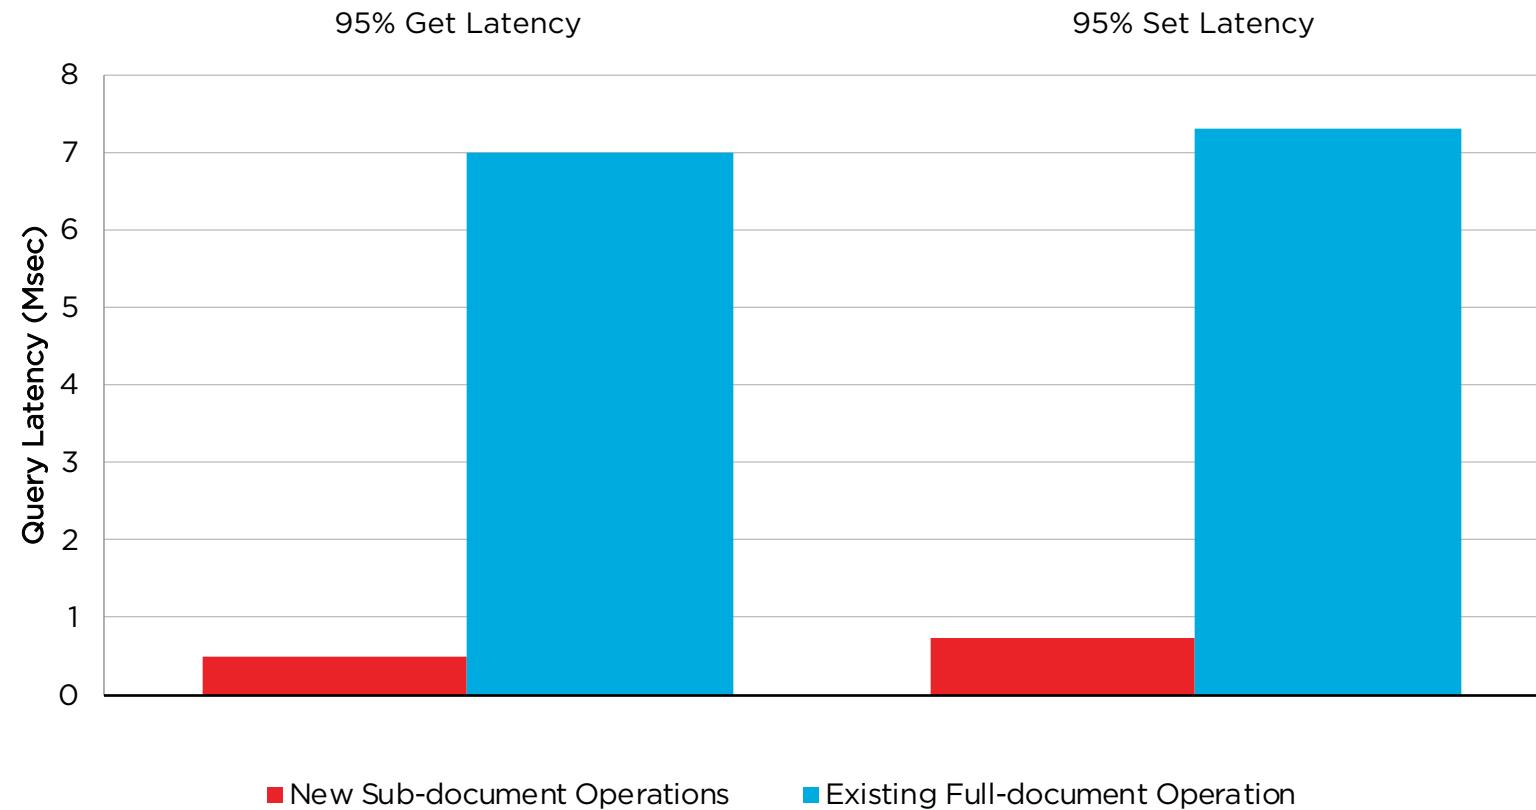




Faster Key Based Operations!

10-14X

Faster Document Read & Update Operation





Nickel (N1QL) : SQL-Like Querying Support

- SQL-like Query Language
 - Expressive, familiar, and feature-rich language for querying, transforming, and manipulating JSON data
 - ANSI 92 SQL Compatible – Selects, Inserts, Updates, Group By, Sort, Functions etc.
- N1QL extends SQL to handle data that is:
 - **Nested:** Contains nested objects, arrays
 - **Heterogeneous:** Schema-optional, non-uniform
 - **Distributed:** Partitioned across a cluster

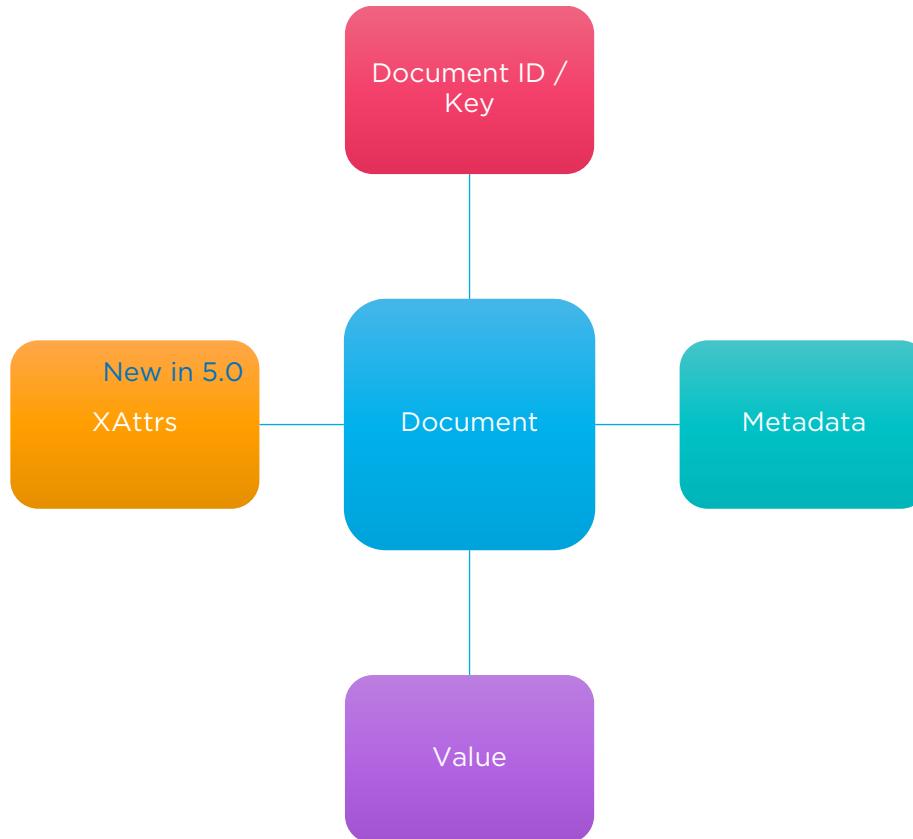


Power of
SQL

Flexibility
of JSON



JSON Document Support



- Document ID / Key (Max 250 bytes):
 - Must be unique / Lookup is extremely fast
 - Similar to primary keys in relational databases
 - Documents are partitioned based on the document ID
- Value (Max 20 MB)
 - JSON
 - Binary - integers, strings, booleans
 - Common binary values include serialized objects, compressed XML, compressed text, encrypted values
- Metadata (Fixed 56 bytes)
 - CAS Value (unique identifier for concurrency)
 - TTL
 - Flags (optional client library metadata)
 - Revision ID #
- XAttr (Max 20 MB)^{New in 5.0}
 - Non-enumerable eXtended Attributes

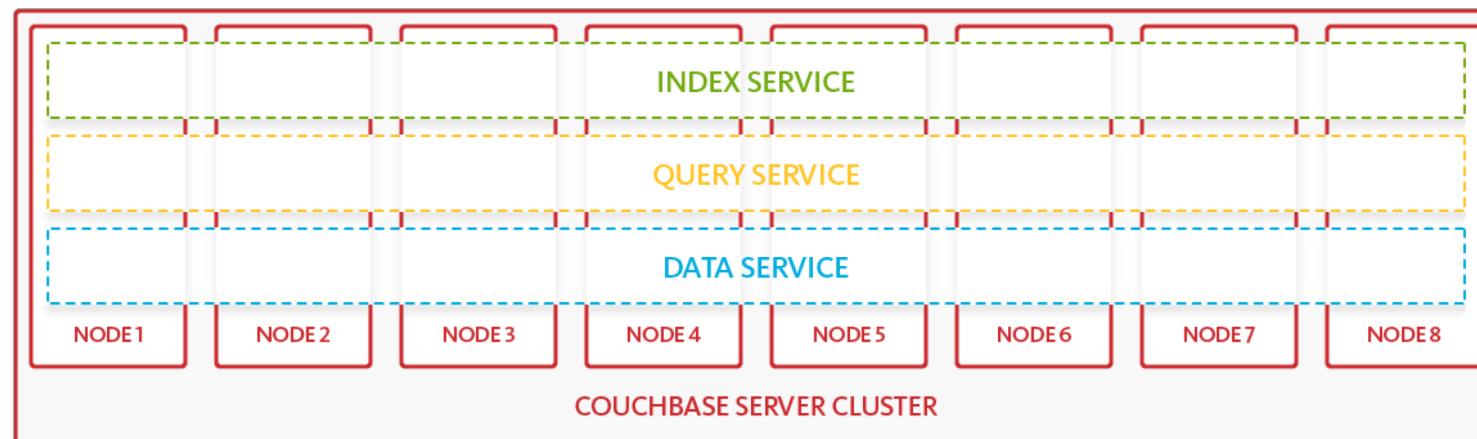


4 MDS



Modern Architecture - Multi-Dimensional Scaling

MDS is the architecture that enables independent scaling of data, query, and indexing workloads while being managed as one cluster.





Modern Architecture - Multi-Dimensional Scaling

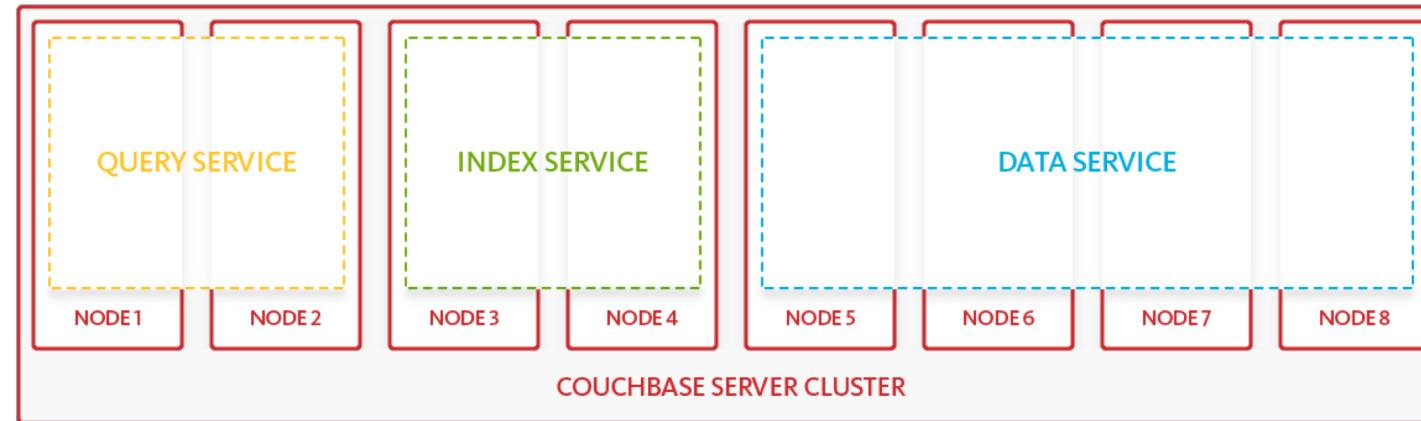
Independent Scalability for Best Computational Capacity — *per Service*.

Heavier Indexing?

Scale up or out
Index Service Nodes.

More RAM for Query Processing?

Scale up or out
Query Service Nodes.





Modern Architecture - Multi-Dimensional Scaling

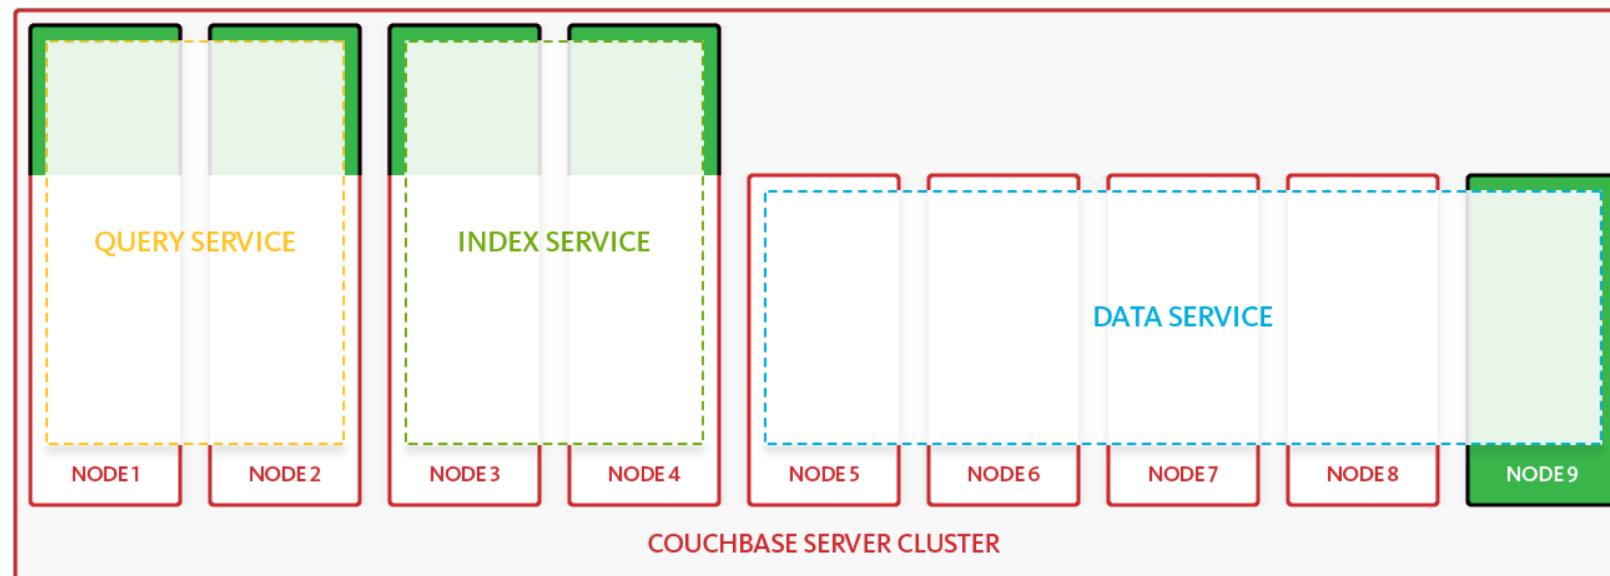
Independent Scalability for Best Computational Capacity — per Service.

Heavier Indexing?

Scale up or out
Index Service Nodes.

More RAM for Query Processing?

Scale up or out
Query Service Nodes.



The same applies for other services



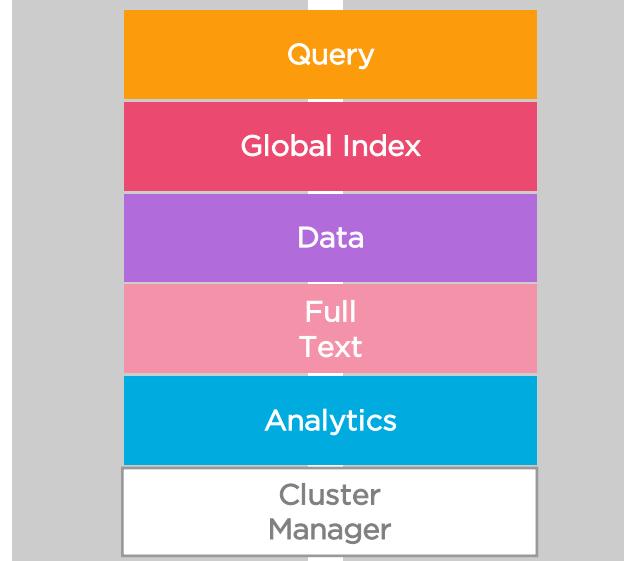
5

Core Principles

#1 Elastic Scaling Architecture



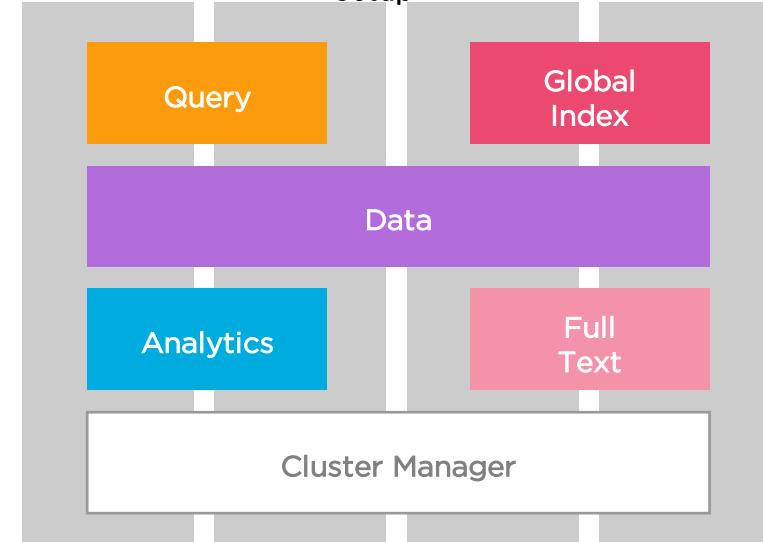
Sample Dev Setup



NODE 1

NODE 2

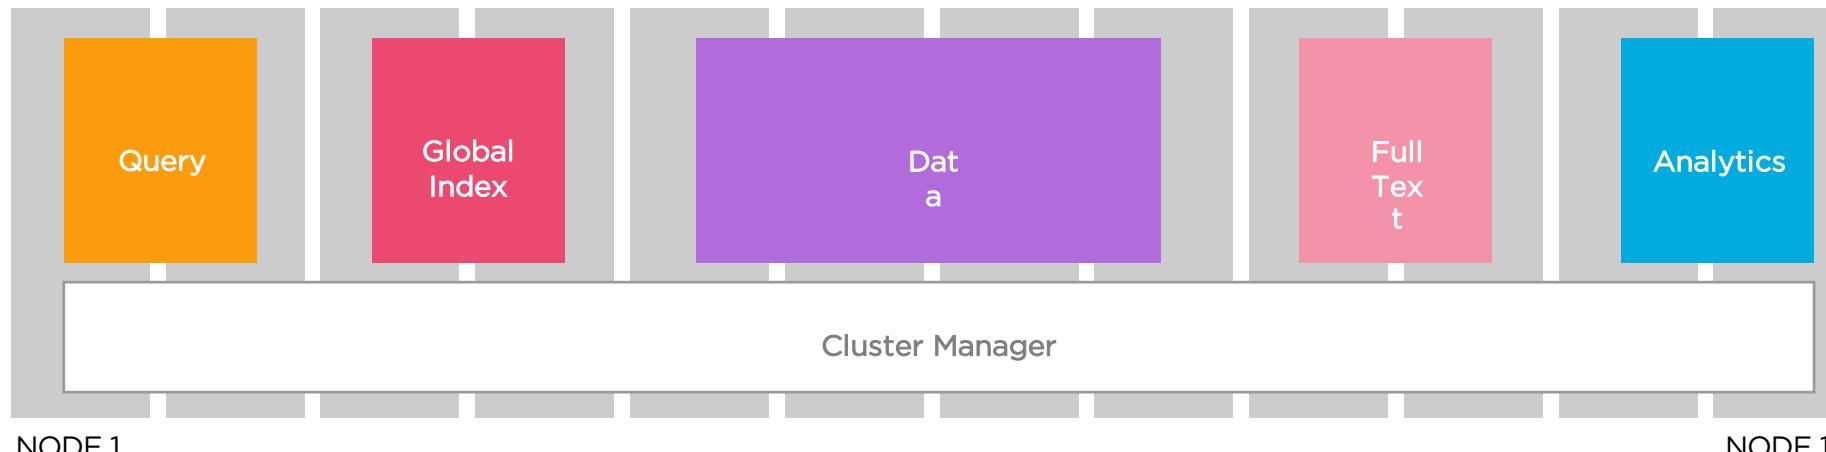
Sample QA Setup



NODE 1

NODE 4

Sample Production Deployment



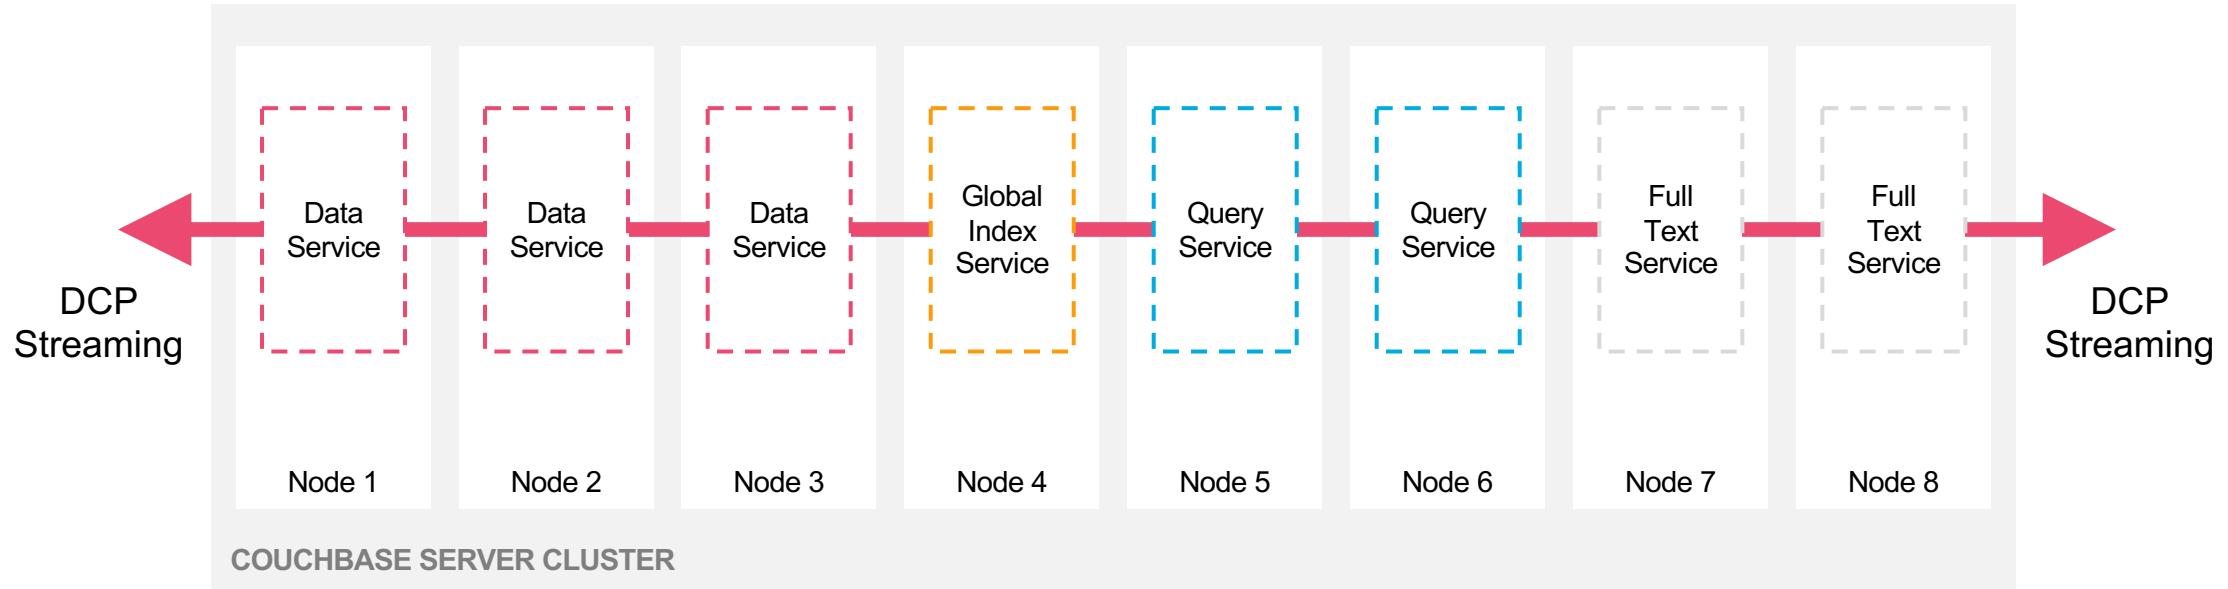
NODE 1

NODE 12

#2 Memory-first architecture



Data movement free from disk bottlenecks



- In-memory streaming of updates to all components
- In-memory cache
- Memory-only data buckets
- Memory-only indexes

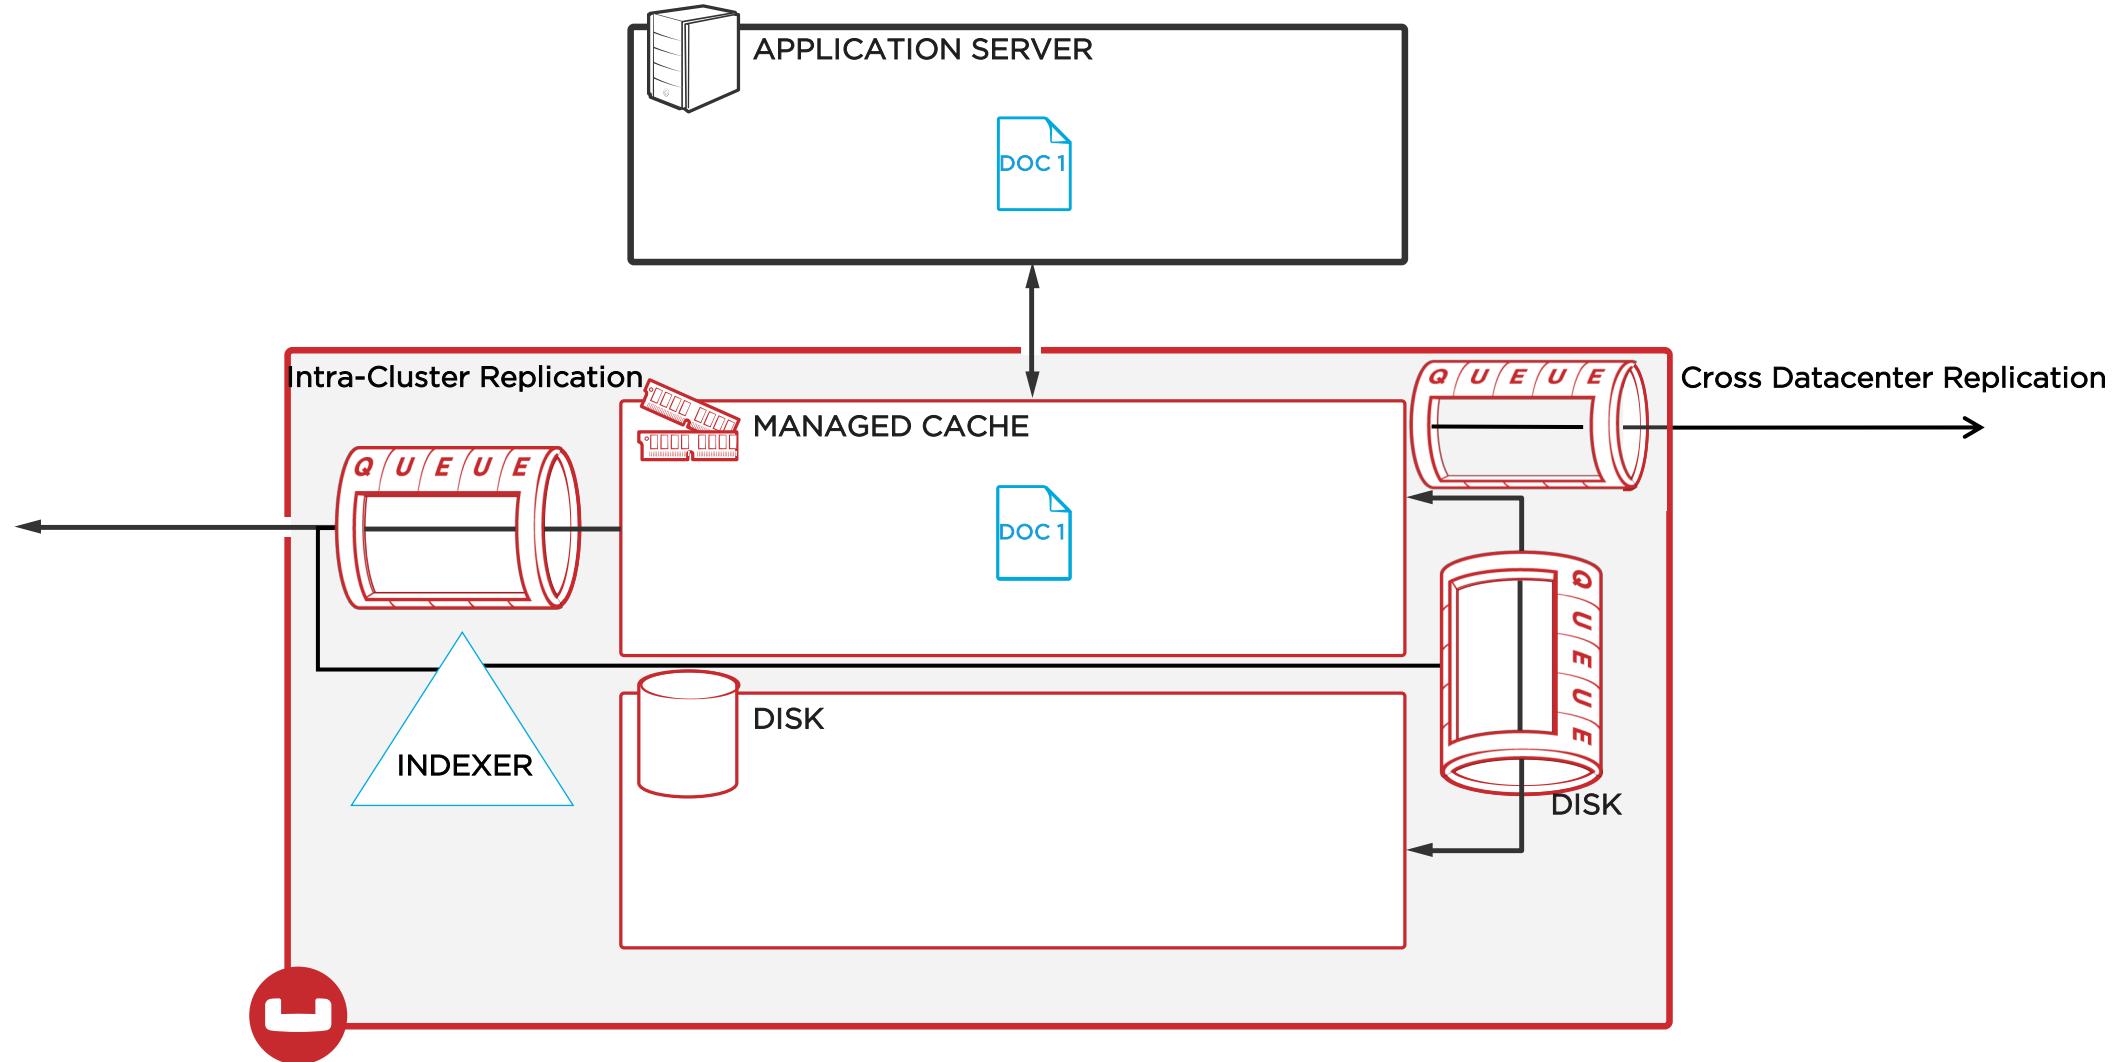
In-memory streaming of updates to all components
In-memory (cached) access to data and indexes
Memory-only indexes

#3 Asynchronous approach to everything



- Persistence
- Intra-cluster Replication
- Inter-cluster Replication
- Global secondary Indexing updates
- Full-Text Search update
- Analytics service updates

#3 Asynchronous approach to everything



#3 Asynchronous approach to everything



Configurable consistency per request / query

Data Consistency

Data access is strongly consistent within cluster
Eventually consistent across clusters

Query Consistency

Specify level of consistency for queries

Thank you



Couchbase